mclust Version 4 for R: Normal Mixture Modeling for Model-Based Clustering, Classification, and Density Estimation

Chris Fraley, Adrian E. Raftery
University of Washington, USA

T. Brendan Murphy
University College Dublin, Ireland

Luca Scrucca
Università degli Studi di Perugia, Italy

mclust is a contributed R package for model-based clustering, classification, and density estimation based on finite normal mixture modeling. It provides functions for parameter estimation via the EM algorithm for normal mixture models with a variety of covariance structures, and functions for simulation from these models. Also included are functions that combine model-based hierarchical clustering, EM for mixture estimation and the Bayesian Information Criterion (BIC) in comprehensive strategies for clustering, density estimation and discriminant analysis. There is additional functionality for displaying and visualizing the models along with clustering, classification, and density estimation results. Several features of the software have been changed in this version, in particular the functionality for discriminant analysis and density estimation has been largely expanded. mclust is licensed under the GPL and distributed through CRAN; see http://cran.r-project.org/web/packages/mclust/index.html.

Contents

List of Tables

List of Figures

1 Overview

The mclust software [12, 14] currently includes the following features:

- Normal mixture modeling via EM for ten parameterized covariance structures;

- Simulation from parameterized Gaussian mixtures;

- Model-based clustering (covariance parameterization and number of clusters selected via BIC);

- Model-based hierarchical clustering for four covariance structures;

- Discriminant analysis based on finite mixture modeling (EDDA, MclustDA);

- Density estimation;

- Displays, including uncertainty plots, random projections, contour and perspective plots, classification plots, density curves in one and two dimensions;

- Combining Gaussian mixture components for clustering;

- Dimension reduction methods for model-based clustering and classification.

This manuscript describes Version 4 of mclust for R, with added functionality for displaying and visualizing the models along with clustering, classification, and density estimation results. A number of other aspects of the software have been changed as well, to reflect evolution in its use, in particular the functionality for discriminant analysis and density estimation has been largely expanded. A comprehensive treatment of the methods used in mclust can be found in [13, 17, 31, 3].

mclust is licensed under the GPL and is available as the contributed package for the R language. It can be obtained from CRAN at http://cran.r-project.org/web/packages/mclust/index.html. Follow the instructions for installing R packages on your machine, and then do

```
> library(mclust)
```

inside the R console to use the software. Throughout this manual it will be assumed that these steps have been taken before running the examples.

Several contributed R packages have functions that require mclust including clustvarsel, fpc, prabclus, msir, and the Bioconductor package spotSegmentation [15]. A couple of tutorials on mclust have also been published [16, 18].

2 Model-Based Cluster Analysis

mclust provides functionality for cluster analysis combining model-based hierarchical clustering (section 5), EM for Gaussian mixture models (section 3), and BIC (section 4).

2.1 Basic Cluster Analysis Example using Mclust

As an illustration, consider the bivariate faithful dataset (included in the R language distribution) shown in Figure 1.

The following command performs a cluster analysis of the faithful dataset, and prints a summary of the results:

Figure 1: The bivariate `faithful` dataset.

```
> faithfulMclust <- Mclust(faithful)
> summary(faithfulMclust)
----------------------------------------------------
Gaussian finite mixture model fitted by EM algorithm
----------------------------------------------------

Mclust EEE (elliposidal, equal volume, shape and orientation) model with 3 components:

 log.likelihood   n df     BIC
         -1126.4 272 11 -2314.4

Clustering table:
   1   2   3
 130  97  45
```

In this case, the best model according to BIC is an equal-covariance model with 3 components or clusters. A more detailed summary including the estimated parameters can be obtained with the following code:

```
> summary(faithfulMclust, parameters = TRUE)
----------------------------------------------------
Gaussian finite mixture model fitted by EM algorithm
----------------------------------------------------

Mclust EEE (elliposidal, equal volume, shape and orientation) model with 3 components:

 log.likelihood   n df     BIC
```

```
        -1126.4 272 11 -2314.4

Clustering table:
  1   2   3
130  97  45

Mixing probabilities:
      1       2       3
0.46190 0.35646 0.18164

Means:
            [,1]    [,2]    [,3]
eruptions  4.4761  2.0378  3.8199
waiting   80.8922 54.4935 77.6711

Variances:
[,,1]
          eruptions waiting
eruptions   0.07728  0.4765
waiting     0.47650 33.7485
[,,2]
          eruptions waiting
eruptions   0.07728  0.4765
waiting     0.47650 33.7485
[,,3]
          eruptions waiting
eruptions   0.07728  0.4765
waiting     0.47650 33.7485
```

The clustering results can be displayed as follows:

```
> plot(faithfulMclust)
```

The corresponding plots are shown in Figure 2.

The covariance structures defining the models available in mclust are summarized in Table 1; these models are explained in more detail in appendix A.1.

The input to function Mclust includes the number of mixture components and the covariance structures to consider. By default, Mclust compares BIC values for parameters optimized for up to nine components and all ten covariance structures currently available in the mclust software. The output includes the parameters of the maximum-BIC model (where the maximum is taken over all of the models and numbers of components considered), and the corresponding classification and uncertainty.

The object produced by Mclust is a list with a components describing the estimated model. The names of these components can be displayed as follows:

```
> names(faithfulMclust)
 [1] "call"          "modelName"     "n"              "d"
 [5] "G"             "BIC"           "bic"            "loglik"
 [9] "df"            "parameters"    "classification" "uncertainty"
[13] "z"
```

and they are described in detail in the Mclust help file.
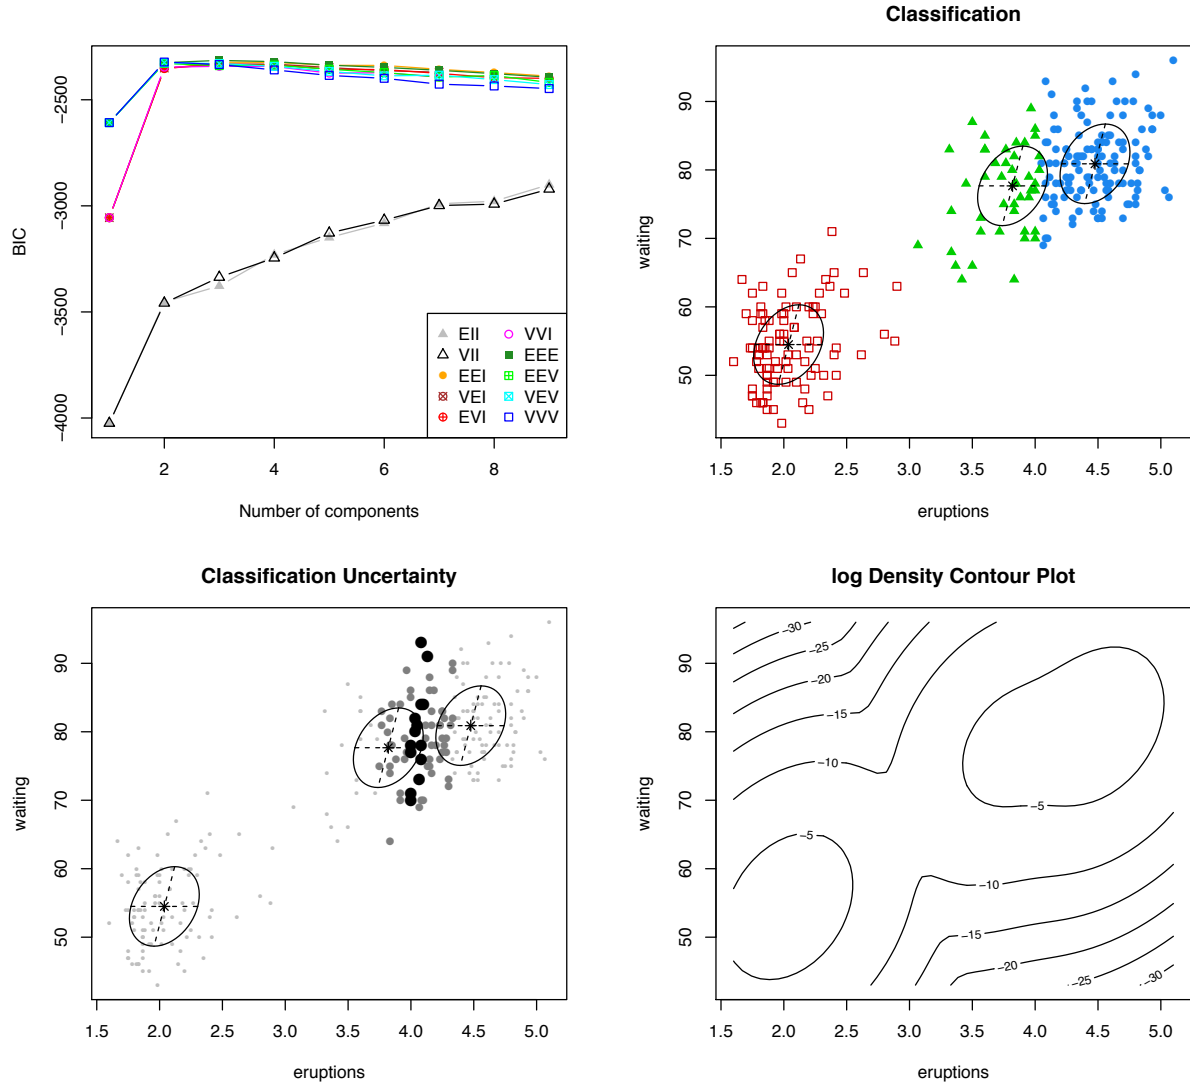
Figure 2: Plots associated with the function `Mclust` for the `faithful` dataset with the default arguments. Clockwise from upper left: BIC, classification, uncertainty, density. The ellipses superimposed on the classification and uncertainty plots correspond to the covariances of the components.

2.2 `mclustBIC` and its `summary` function

To do further analysis on the same dataset, for example to see the results for a different set of models and/or different numbers of components, `Mclust` could be rerun. However this approach could involve unnecessary repetition of computations and could also take considerable time when the dataset is large or the process is to be repeated many times. An alternative approach is to split the analysis into several parts using function `mclustBIC`.

For the `faithful` dataset, the following sequence of commands produces the same clustering result as the call to `Mclust`.

```
> faithfulBIC <- mclustBIC(faithful)
> faithfulSummary <- summary(faithfulBIC, data = faithful)
> faithfulSummary
```

Table 1: Parameterizations of the covariance matrix Σ_k currently available in mclust for hierarchical clustering (HC) and/or EM for multidimensional data. ('•' indicates availability).

identifier	Model	HC	EM	Distribution	Volume	Shape	Orientation
E		•	•	(univariate)	equal		
V		•	•	(univariate)	variable		
EII	λI	•	•	Spherical	equal	equal	NA
VII	$\lambda_k\,I$	•	•	Spherical	variable	equal	NA
EEI	λA		•	Diagonal	equal	equal	coordinate axes
VEI	$\lambda_k A$		•	Diagonal	variable	equal	coordinate axes
EVI	λA_k		•	Diagonal	equal	variable	coordinate axes
VVI	$\lambda_k A_k$		•	Diagonal	variable	variable	coordinate axes
EEE	$\lambda D A D^T$	•	•	Ellipsoidal	equal	equal	equal
EEV	$\lambda D_k A D_k^T$		•	Ellipsoidal	equal	equal	variable
VEV	$\lambda_k D_k A D_k^T$		•	Ellipsoidal	variable	equal	variable
VVV	$\lambda_k D_k A_k D_k^T$	•	•	Ellipsoidal	variable	variable	variable

```
classification table:
  1   2   3
130  97  45

best BIC values:
  EEE,3   EEE,4   VVV,2
-2314.4 -2320.2 -2322.2
```

Although the method used for printing is different,
faithfulSummary is nearly identical to the object faithfulMclust computed in the previous section, the difference being that it does not include the "BIC" component giving the table of BIC values, which comprise the object faithfulBIC computed by mclustBIC:

```
> faithfulBIC
 BIC:
      EII     VII     EEI     VEI     EVI     VVI     EEE     EEV     VEV     VVV
1 -4024.7 -4024.7 -3055.8 -3055.8 -3055.8 -3055.8 -2607.6 -2607.6 -2607.6 -2607.6
2 -3453.0 -3458.3 -2354.6 -2350.6 -2352.6 -2346.1 -2325.2 -2329.1 -2325.4 -2322.2
3 -3377.7 -3336.5 -2323.0 -2332.7 -2332.2 -2342.4 -2314.4 -2339.0 -2329.4 -2333.9
4 -3230.2 -3245.7 -2323.7 -2331.8 -2334.8 -2343.1 -2320.2 -2336.8 -2342.5 -2359.2
5 -3149.4 -3128.2 -2337.7 -2348.3 -2355.9 -2374.3 -2337.0 -2356.2 -2366.2 -2385.3
6 -3081.4 -3067.6 -2338.1 -2363.1 -2357.7 -2372.7 -2347.3 -2371.7 -2387.4 -2399.0
7 -2990.3 -2998.5 -2356.5 -2370.1 -2375.9 -2393.1 -2361.2 -2393.0 -2384.2 -2426.5
8 -2978.1 -2991.9 -2371.8      NA -2396.0      NA -2376.9 -2385.8 -2404.9 -2435.0
9 -2899.8 -2921.0 -2388.6      NA -2399.1      NA -2393.7 -2418.3 -2428.4 -2447.3
```

The missing values correspond to models and numbers of clusters for which parameter values could not be fit (using the default initialization). For multivariate data, the default initialization for all models uses the classification from hierarchical clustering based on an unconstrained model. For univariate data, the default is to classify the data by quantile for initialization.

The summary method for mclustBIC allows specification of the models and numbers of clusters over which the best model is to be chosen, allowing models other than the maximum BIC model

to be extracted and analyzed.

The `plot` method for `mclustBIC` allows specification of the models and numbers of clusters, arguments to the `legend` function, as well as setting limits on the vertical axis. For example, the following shows the maximum BIC values in more detail than the default:

```
> plot(faithfulBIC, G = 1:7, ylim = c(-2500,-2300), legendArgs = list(x = "bottomright", ncol = 5))
```

The resulting plot is shown in Figure 3.
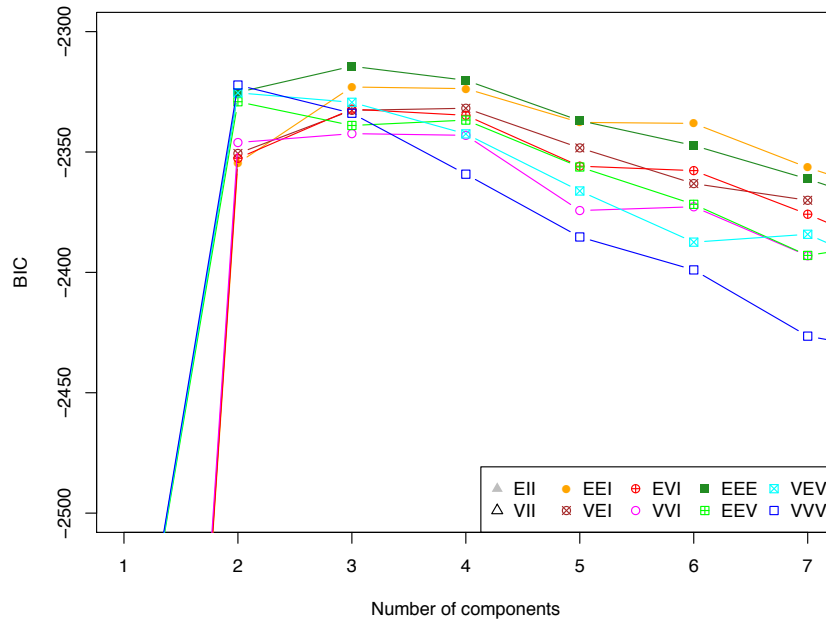
Figure 3: BIC plot for the `faithful` dataset, with vertical axes adjusted to display the maximum values.

2.3 Extended Cluster Analysis Example

As an example of an extended analysis, consider the `wreath` data shown in Figure 4. There are 1000 bivariate observations simulated from a 14-component model in which the component covariance matrices are of equal size and shape, but differ in orientation.

```
> data(wreath)
> plot(wreath, pch = 20, cex = 0.5)
```

The BIC values can be obtained with a call to `mclustBIC` and then plotted:

```
> wreathBIC <- mclustBIC(wreath)
> plot(wreathBIC, legendArgs = list(x = "topleft"))
```

Refering to the BIC plot (shown on the left in Figure 5), the maximum BIC appears to be outside the range of the default values for the number of components in `mclustBIC` (and `Mclust`).

More components (for example, up to 20) can be considered in the analysis without recomputing previous results:

```
> wreathBIC <- mclustBIC(wreath, G = 1:20, x = wreathBIC)
> plot(wreathBIC, G = 10:20, legendArgs = list(x = "bottomright"))
> summary(wreathBIC, wreath)
```

9

Figure 4: The bivariate `wreath` dataset, which consists of 1000 observations simulated from a 14-component normal mixture in which the component covariance matrices are of equal size and shape, but differ in orientation.

```
classification table:
 1  2  3  4  5  6  7  8  9 10 11 12 13 14
74 69 63 74 68 70 71 66 83 77 66 77 61 81

best BIC values:
EEV,14 EEV,15 EEV,16
-10903 -10926 -10953
```

The BIC plot is shown on the right in Figure 5. Using `summary` to obtain the best model according to BIC, a 14-component `EEV` model is chosen, which is in agreement with how the data was simulated.

```
> wreathModel <- summary(wreathBIC, data = wreath)
> wreathModel
classification table:
 1  2  3  4  5  6  7  8  9 10 11 12 13 14
74 69 63 74 68 70 71 66 83 77 66 77 61 81

best BIC values:
EEV,14 EEV,15 EEV,16
-10903 -10926 -10953
```

The model for the `wreath` dataset is shown in Figure 6.

The `summary` function can also be used to restrict the set of models and/or numbers of clusters over which the best model is chosen according to BIC. For example, the following commands
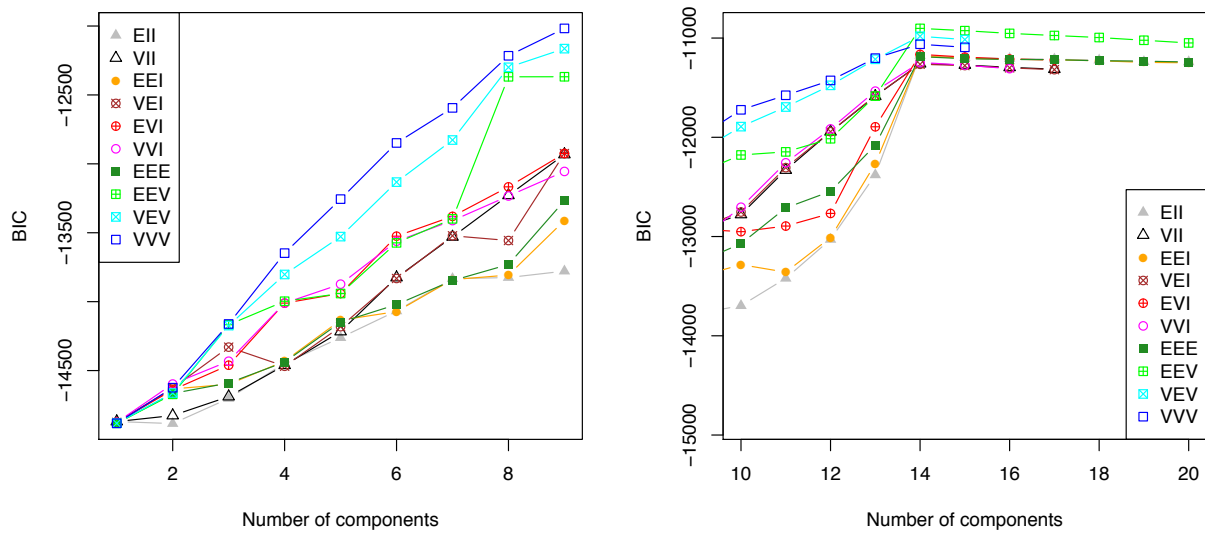
Figure 5: BIC for `wreath` dataset. LEFT: BIC for all models and up to 9 components (the default in `mclustBIC` and `Mclust`). RIGHT: BIC for 10:20 components, all models. There is a clear peak for all models at 14 components.

produce the best spherical model for the `wreath` data:

```
> wreathSphericalModel <- summary(wreathBIC, data = wreath,
                                  modelNames = c("EII", "VII"))
> wreathSphericalModel
classification table:
 1  2  3  4  5  6  7  8  9 10 11 12 13 14
75 69 63 74 68 70 71 65 83 77 66 77 61 81

best BIC values:
EII,14 EII,15 EII,16
-11176 -11197 -11207
```

2.4   Clustering Univariate Data

Cluster analysis for univariate data can be carried out as for two and higher dimensions. As an example, we use the `precip` dataset (included in the R language distribution):

```
> precipMclust <- Mclust(precip)
> summary(precipMclust, parameters = TRUE)
----------------------------------------------------
Gaussian finite mixture model fitted by EM algorithm
----------------------------------------------------

Mclust V (univariate, unequal variance) model with 2 components:

 log.likelihood  n df     BIC
        -275.47 70  5 -572.19

Clustering table:
```

11

Figure 6: The 14-component `EEV` (equal size and shape) model obtained for the `wreath` dataset. The ellipses superimposed on the plot correspond to the covariances of the components.

```
 1  2
13 57

Mixing probabilities:
      1        2
0.18181 0.81819

Means:
      1        2
12.806 39.792

Variances:
      1        2
16.891 90.189

> plot(precipMclust, legendArgs = list(x = "bottomleft"))
```

Figure 7 shows the BIC, classification, uncertainty, and density for this example.

2.5   Regularizing with a Prior

`mclust` allows to specify a prior distribution to regularize the fit to the data [17]. We illustrate the use of a prior on the `trees` dataset (included in the `R` language distribution), for which a pairs plot
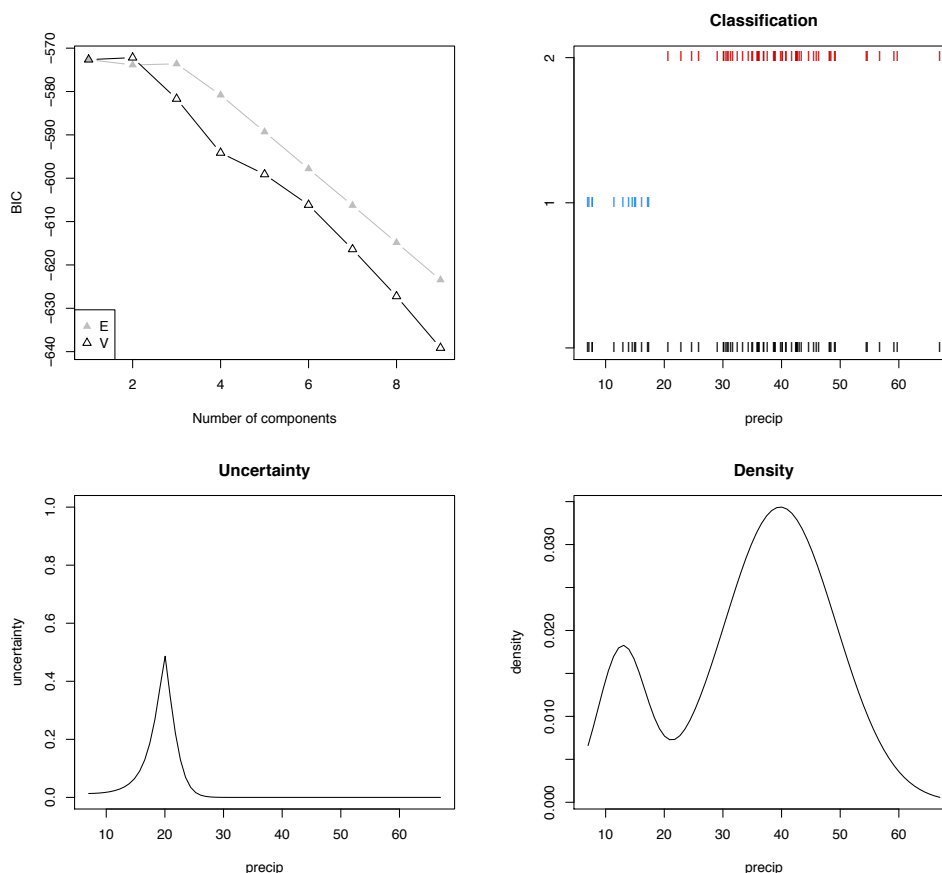
Figure 7: Model-based clustering of the univariate R dataset `precip`. Clockwise from upper left: BIC, classification, uncertainty, and density from `Mclust` applied to the simulated univariate example. In the classification plot, all of the data is displayed at the bottom, with the separated classes shown on different levels above.

is shown in Figure 8.

The following commands compute and plot the BIC curves for the `trees` dataset provided in R with and without a `prior`:

```
> treesBIC <- mclustBIC(trees) # default (no prior)
> plot(treesBIC, legendArgs = list(x = "bottom", ncol = 2, cex = .75))
> treesBICprior <- mclustBIC(trees, prior = priorControl())
> plot(treesBICprior, legendArgs = list(x = "bottom", ncol = 2, cex = .75))
```

Without the prior, the BIC plot shows a number of jagged peaks, and many BIC values are missing for some models due to failure in the EM computations caused by singularity and/or shrinking components. With the prior, the BICs are smoother and there are fewer failures in estimation. See Figure 9.

A function `priorControl` is provided in `mclust` for specifying the prior and its parameters. When called with its defaults, it invokes another function called `defaultPrior` which can serve as a template for specifying alternative priors. An example of the result of a call to `defaultPrior` is shown below.

```
> defaultPrior(trees, G=2, modelName = "VVV")
```

13
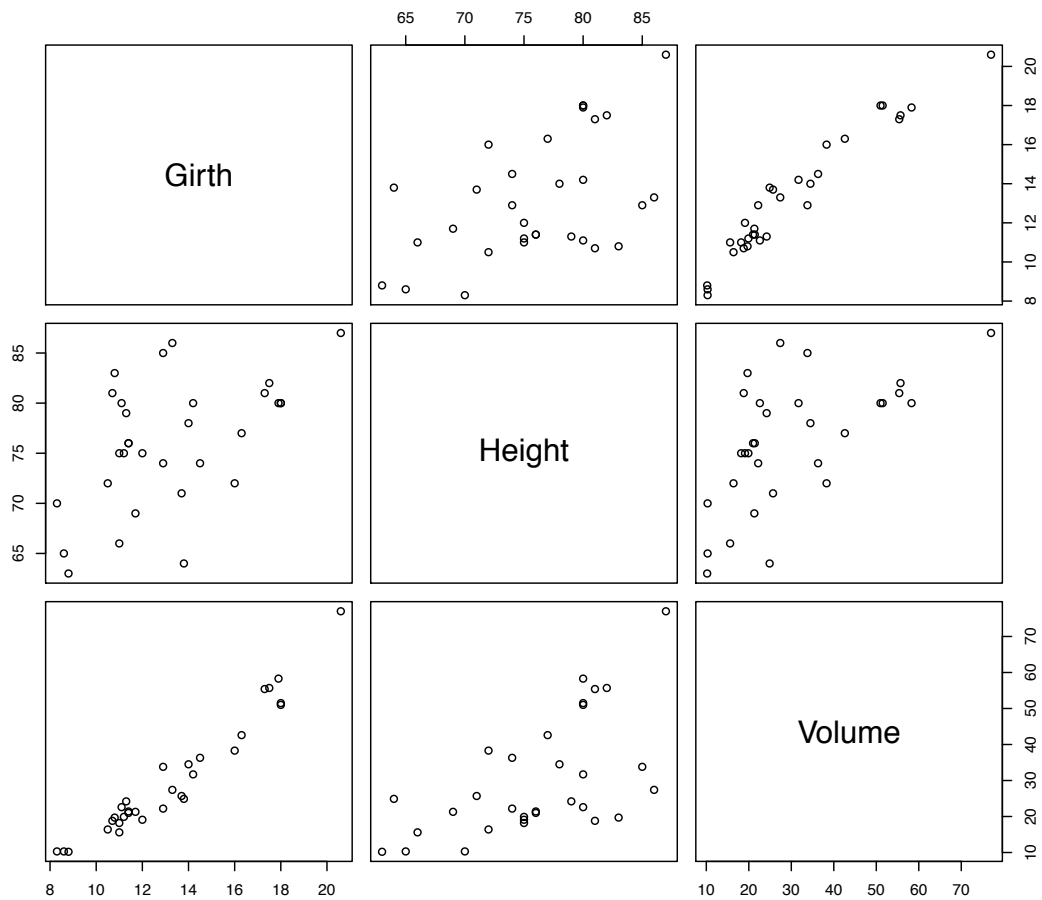
Figure 8: Pairs plot of `trees` dataset.

```
$shrinkage
[1] 0.01

$mean
 Girth Height Volume
13.248 76.000 30.171

$dof
[1] 5

$scale
        Girth  Height  Volume
Girth   6.2038  6.5411  31.428
Height  6.5411 25.5764  39.473
Volume 31.4275 39.4733 170.217
```

For more detail on the prior and its specification, see Section A.4.
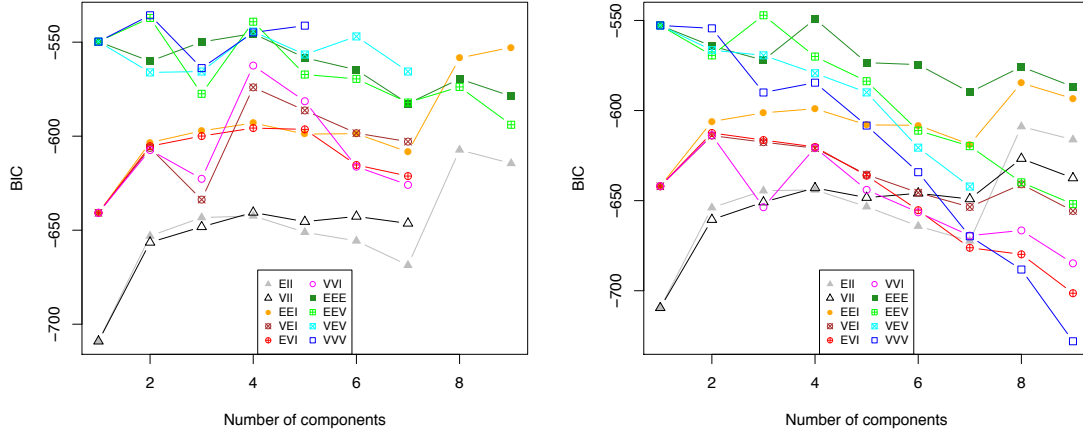
Figure 9: BIC without (left) and with the prior (right) for the `trees` dataset.

2.6   Clustering with Noise and Outliers

mclust allows model-based clustering with noise, namely outlying observations that do not belong to any cluster. To include noise in the modeling, an initial guess of the noise observations must be supplied via the `noise` component of the `initialization` argument in `Mclust` or `mclustBIC`. The model for noise used in mclust is discussed in more detail in Section A.2 of the appendix, along with some some strategies for obtaining an initial noise estimate.

In the following example, Poisson noise is added to the `faithful` dataset. A random initial estimate was used for noise for the purposes of illustration. This happens to work well in this instance, although we don't recommend this as a general strategy.

```
> b <- apply( faithful, 2, range)
> nNoise <- 500
> set.seed(0)
> poissonNoise <- apply(b, 2, function(x, n)
                        runif(n, min = min(x)-.1, max = max(x)+.1), n = nNoise)
> faithfulNdata <- rbind(faithful, poissonNoise)
> set.seed(0)
> faithfulNoiseInit <- sample(c(TRUE,FALSE),size=nrow(faithful)+nNoise,
                              replace=TRUE,prob=c(3,1))
> faithfulNbic <- mclustBIC(faithfulNdata,
                            initialization = list(noise = faithfulNoiseInit))
> faithfulNsummary <- summary(faithfulNbic, faithfulNdata)
> faithfulNsummary
classification table:
  0   1   2
532 141  99

best BIC values:
  EVI,2   VVI,2   EEI,2
-7982.0 -7982.9 -7992.8
```

The results are shown in Figure 10. The classification and BIC plots were obtained with the following commands.

15

```
> plot(faithful)
> points(poissonNoise, pch = 20, cex = 0.3, col = "lightgrey")
> mclust2Dplot(faithfulNdata, classification=faithfulNsummary$classification,
                parameters=faithfulNsummary$parameters)
> plot(faithfulNbic, legendArgs = list(x = "bottomleft", horiz = TRUE, cex = 0.75))
```
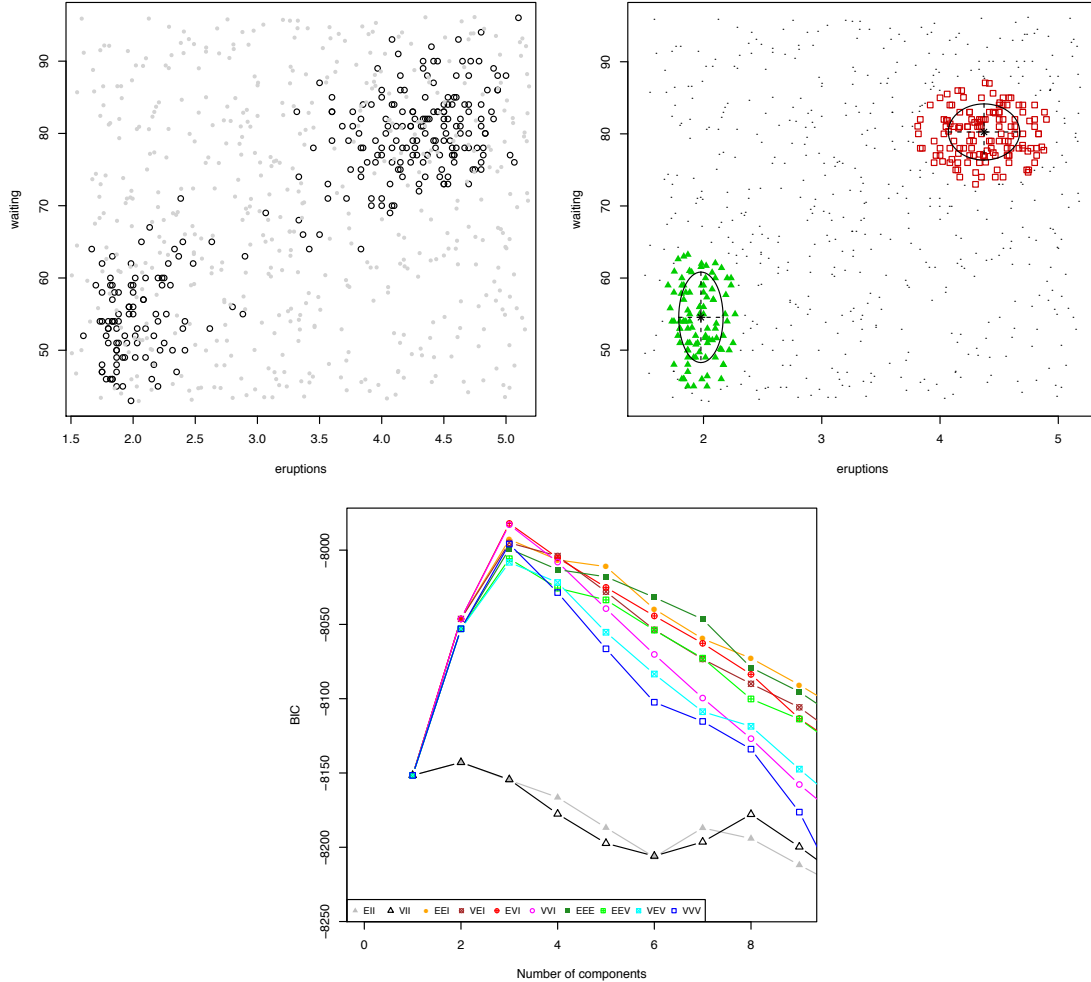

Figure 10: Cluster analysis of the `faithful` dataset with added Poisson noise. Upper Left: The 272 observations of the `faithful` dataset (circles) with 500 Poisson noise points (small dots). Upper Right: `mclust` classification starting with random noise estimate. Lower: BIC.

2.7 Further Considerations in Cluster Analysis

Clustering can be affected by parameter settings such as convergence tolerances within the clustering functions, although the defaults are often adequate. It is also possible do model-based clustering starting with parameter estimates, conditional probabilities, or classifications other than those produced by model-based hierarchical clustering. The functions provided for mixture estimation (Section 3) and BIC (Section 4) can be used for this purpose.

Finally, it is important to take into account numerical issues in cluster analysis. The computations for estimating the model parameters break down when the covariance corresponding to one or

more components becomes ill-conditioned (singular or nearly singular). Including a prior (Section A.4) is often helpful in such situations. In general the modeling computations cannot proceed if clusters contain only a few observations or if the observations they contain are very nearly colinear. Computations may also fail when one or more mixing proportions shrink to negligible values. The EM functions in mclust compute and monitor the conditioning of the covariances, and an error condition is issued (unless such warnings are turned off) when the associated covariance appears to be nearly singular, as determined by a threshold with the default value `emControl()$eps`.

3 EM for Mixture Models

mclust provides iterative EM (Expectation-Maximization) methods for maximum-likelihood estimation in parameterized Gaussian mixture models. In the models considered here, an iteration of EM consists of an 'E'-step, which computes a matrix z such that z_{ik} is an estimate of the conditional probability that observation i belongs to group k given the current parameter estimates, and an 'M'-step, which computes parameter estimates given z.

mclust functions `em` and `me` implement the EM algorithm for parameterized Gaussian mixtures. Function `em` starts with the E-step; besides the data and model specification, the model parameters (means, covariances, and mixing proportions) proportions must be provided. Function `me` starts with the M-step; besides the data and model specification, the conditional probabilities z must be provided. The output for both are the maximum-likelihood estimates of the model parameters and z.

3.1 Individual E and M Steps

Functions `estep` and `mstep` implement the individual steps of the EM iteration. Conditional probabilities z and the log likelihood can be recovered from parameters via `estep`, while parameters can be recovered from conditional probabilities z using `mstep`. Below we apply `mstep` and `estep` to the `iris` dataset(included in the R language distribution).

```
> ms <- mstep(modelName = "VVV", data = iris[,-5], z = unmap(iris[,5]))
> names(ms)
[1] "modelName"  "prior"      "n"          "d"          "G"          "z"
[7] "parameters"
> es <- estep(modelName = "VVV", data = iris[,-5], parameters = ms$parameters)
> names(es)
[1] "modelName"  "n"          "d"          "G"          "z"          "parameters"
[7] "loglik"
```

In this example, the initial estimate of z for the M-step is a matrix of indicator variables corresponding to a discrete classification (`iris[,5]`). The function `unmap` converts a discrete classification into the corresponding indicator variables. mclust allows specification of a prior, for which the EM algorithm will compute a posterior mode. See Sections 2.5 and A.4 for more details.

3.2 Uncertainty

The uncertainty in the classification associated with conditional probabilities z can be obtained by subtracting the probability of the most likely group for each observation from 1:

```
> meVVViris <- me(modelName = "VVV", data = iris[,-5], z = unmap(iris[,5]))
> uncer <- 1 - apply( meVVViris$z, 1, max)
```

The R function `quantile` applied to the uncertainty gives a measure of the quality of the classification:

```
> quantile(uncer)
        0%         25%         50%         75%        100%
0.0000e+00 0.0000e+00 1.9070e-08 1.3921e-03 3.3619e-01
```

In this case the indication is that the majority of observations are well classified. Note, however, that when groups intersect, uncertain classifications would be expected in the overlapping regions.

When a true classification is known, the relative uncertainty of misclassified observations can be displayed by function `uncerPlot`, as is done below for the `iris` example (see Figure 11):

```
> uncerPlot(z = meVVViris$z, truth = iris[,5])
```

It is also possible to plot an uncertainty curve for one-dimensional data (see Section 2.4) or an uncertainty surface for bivariate data (see Section 8.1).

observations in order of increasing uncertainty
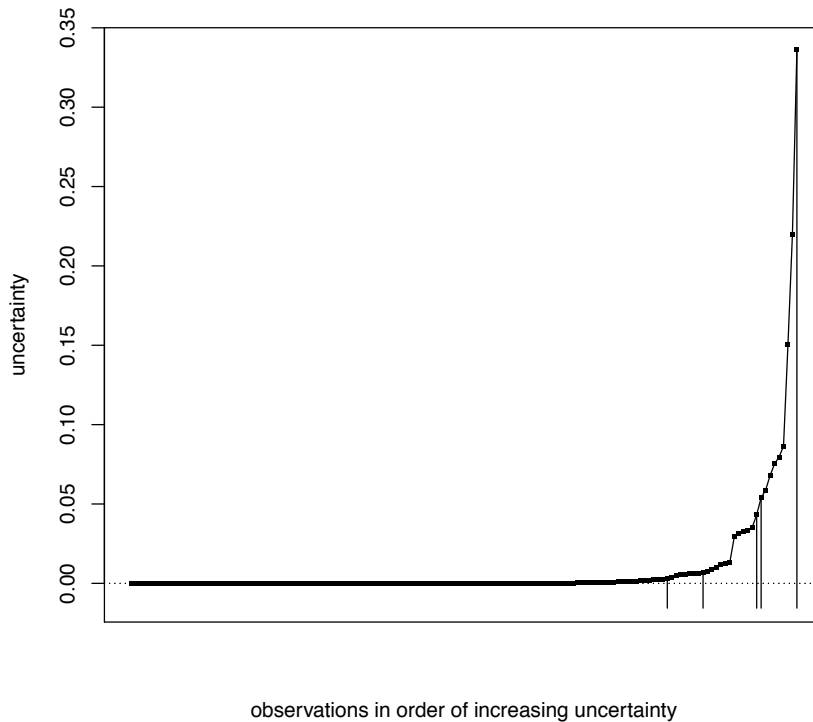
Figure 11: Uncertainty plot for the 3-cluster mixture model fit of the `iris` dataset via EM based on unconstrained Gaussian mixtures. The vertical lines indicate misclassified observations. The plot was created with function `uncerPlot`, and shows the relative uncertainty of misclassified observations.

3.3   Control Parameters

Besides the initial values and the prior, other parameters can influence the outcome of `em` or `me`. These include:

`tol` Iteration convergence tolerance. The default is `emControl()$tol=c(1.e-5,`$\sqrt{\epsilon_M}$`)`, where ϵ_M is the relative machine precision, which has the value `2.220446e-16` on IEEE compliant

machines. The first value is the tolerance for relative convergence of the loglikelihood in the EM algorithm, and the second value is the relative parameter convergence tolerance for the M-step for those models that have an iterative M-step ("VEI", "VEV").

eps A tolerance for terminating iterations due to ill-conditioning, such as near singularity in covariance matrices. The default is emControl()$eps which is set to the relative machine precision ϵ_M.

A function emControl is provided in mclust for setting these parameters and supplying default values. Although these control settings are in a sense hidden by the defaults, they may have a significant effect on results in some instances and should be taken into consideration in analysis.

4 Bayesian Information Criterion

mclust provides a function bic to compute the Bayesian Information Criterion (BIC) [29] given the maximized loglikelihood for model, the data dimensions, and the number of components in the model. The BIC is the value of the maximized log-likelihood with a penalty on the number of model parameters, and allows comparison of models with differing parameterizations and/or differing numbers of clusters. In general the larger the value of the BIC, the stronger the evidence for the model and number of clusters (see, e.g. [13]). The following shows the BIC calculation in mclust for the 3-cluster classification the iris dataset with the unconstrained variance model:

```
> meVVViris <- me(modelName = "VVV", data = iris[,-5], z = unmap(iris[,5]))
> bic(modelName = "VVV", loglik = meVVViris$loglik,
      n = nrow(iris), d = ncol(iris[,-5]), G = 3)
[1] -580.84
```

5 Model-Based Hierarchical Clustering

mclust provides functions hc for model-based hierarchical agglomeration, and hclass for determining the resulting classifications. Function hc implements fast methods based on the multivariate normal classification likelihood [10]. We use the iris dataset distributed with R in our example. Figure 12 is a pairs plot of the iris dataset in which the three species are differentiated by symbol, obtained by the following command:

```
> clPairs(data = iris[,-5], classification = iris[,5])
```

Below we apply the hierarchical clustering algorithm for unconstrained covariances (VVV) to the iris dataset:

```
> hcVVViris <- hc(modelName = "VVV", data = iris[,-5])
```

The classification produced by hc for various numbers of clusters can be obtained with hclass. For example, for the classifications corresponding to 2 and 3 clusters:

```
> cl <- hclass(hcVVViris, 2:3)
```

The classifications can be displayed with the data using clPairs:

```
> clPairs(data = iris[,-5], classification = cl[,"2"])
> clPairs(data = iris[,-5], classification = cl[,"3"])
```
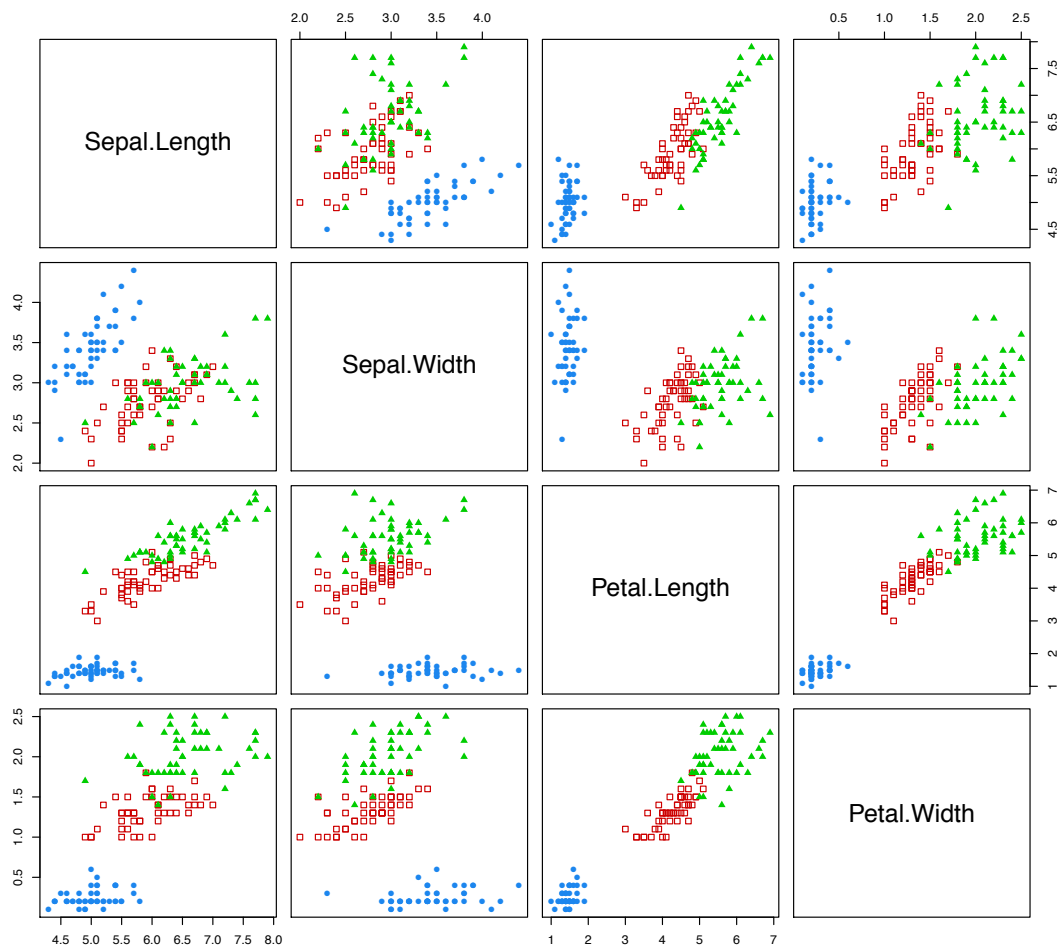
Figure 12: Pairs plot of the `iris` dataset showing classification into species.

More options for displaying clustering and classification results are discussed in Section 8. The 3-group classification can be compared with the known 3-group classification into species, which is given in the 5th column of the `iris` data, using function `classError`:

```
> classError(cl[,"3"], truth = iris[,5])
$misclassified
 [1] 102 107 114 115 120 122 124 127 128 134 139 143 147 150

$errorRate
[1] 0.093333
```

Function `hc` starts by default with every observation of the data in a cluster by itself, and continues until all observations are merged into a single cluster. Arguments `partition` and `minclus` can be used to initialize the process at a chosen nontrivial partition, and to stop it before it reaches the final stage of merging.

Function `hc` for model-based hierarchical clustering based on the unconstrained (`VVV`) model is used to obtain the default initial values for the model-based clustering functions `Mclust` and `mclustBIC`.

6 Density Estimation

The finite mixture estimation capabilities of `mclust` can also be viewed as a general strategy for density estimation.

Consider the following one-dimensional case:

```
> densWaiting <- densityMclust(faithful$waiting)
> summary(densWaiting, parameters = TRUE)
-------------------------------------------------------
Density estimation via Gaussian finite mixture modeling
-------------------------------------------------------

Mclust E (univariate, equal variance) model with 2 components:

 log.likelihood   n df     BIC
          -1034 272  4 -2090.4

Clustering table:
  1   2
 99 173

Mixing probabilities:
      1       2
0.36102 0.63898

Means:
     1      2
54.619 80.094

Variances:
     1      2
34.439 34.439
```

A two-components mixture of Gaussian variables with the same variance is selected by BIC. The parameter estimates can be read from the `summary` output.

A plot of density estimate can be obtained using the corresponding `plot` method:

```
> plot(densWaiting)
```

The density can also be plotted together with a histogram of the observed data by using the optional argument `data`:

```
> plot(densWaiting, data = faithful$waiting)
```

In the one-dimensional case two diagnostic plots [24, pp. 87–90] for density estimation are available. The first plot is a graph of the estimated CDF vs the empirical distribution function (see left panel of Figure 14). The estimated CDF is computed by the function `cdfMclust`, which numerically integrate the density estimate over a grid of points. The second graph is a Q-Q plot of the sample quantiles vs the quantiles computed from the estimated density (see right panel of Figure 14). Both graphs can be obtained as follows:

```
> plot(densWaiting, what = "diagnostic")
```

A plot of BIC values as discussed in Section 2 is also available by setting `what = "BIC"`.

Moving to higher dimensions, we may consider the bivariate `faithful` dataset:
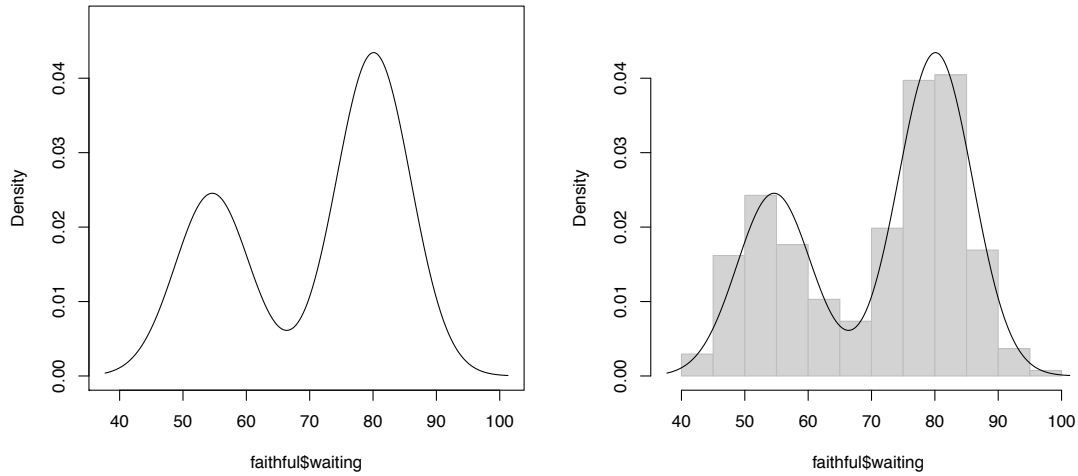
Figure 13: Density estimate (left) and density estimate with histogram of data (right) for the waiting time to next eruption from `faithful` dataset.
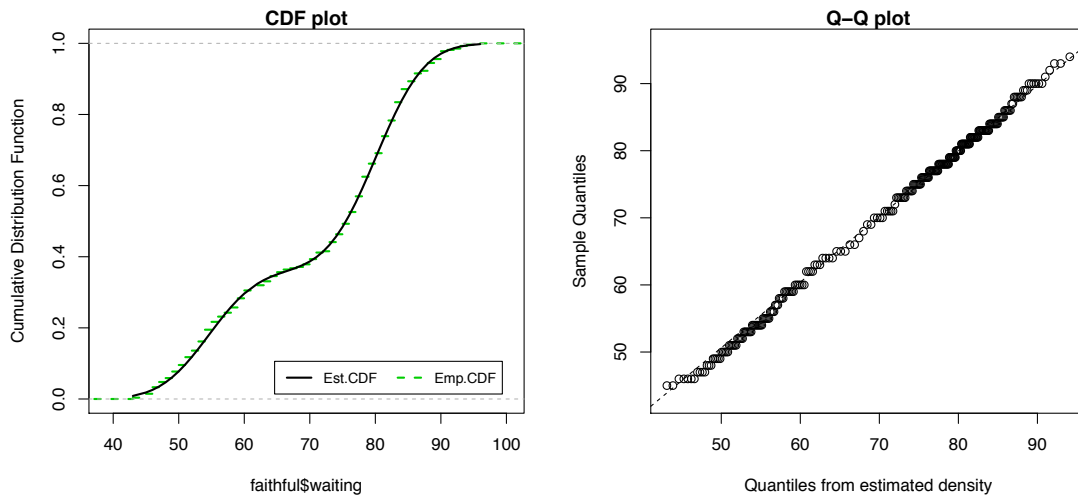
Figure 14: Density estimate diagnostics: plot of estimated CDF and empirical distribution function (left); Q-Q plot of sample quantiles vs quantiles from density estimation.

```
> faithfulDens <- densityMclust(faithful)
```

For two-dimensional data the default plot produces a contour plot of density estimates:

```
> plot(faithfulDens)
> plot(faithfulDens, faithful, col = "grey", nlevels = 10, drawlabels = FALSE)
```

The two plots are shown in the top panels of Figure 15. In the second contour plot we added the sample data and we customize some parameters. A bivariate density estimate may also be plotted using an image plot or a perspective plot as follows:

```
> plot(faithfulDens, type = "image", col = topo.colors(50))
> plot(faithfulDens, type = "persp", col = grey(0.8))
```

The graphs are reported in the bottom panels of Figure 15.

22
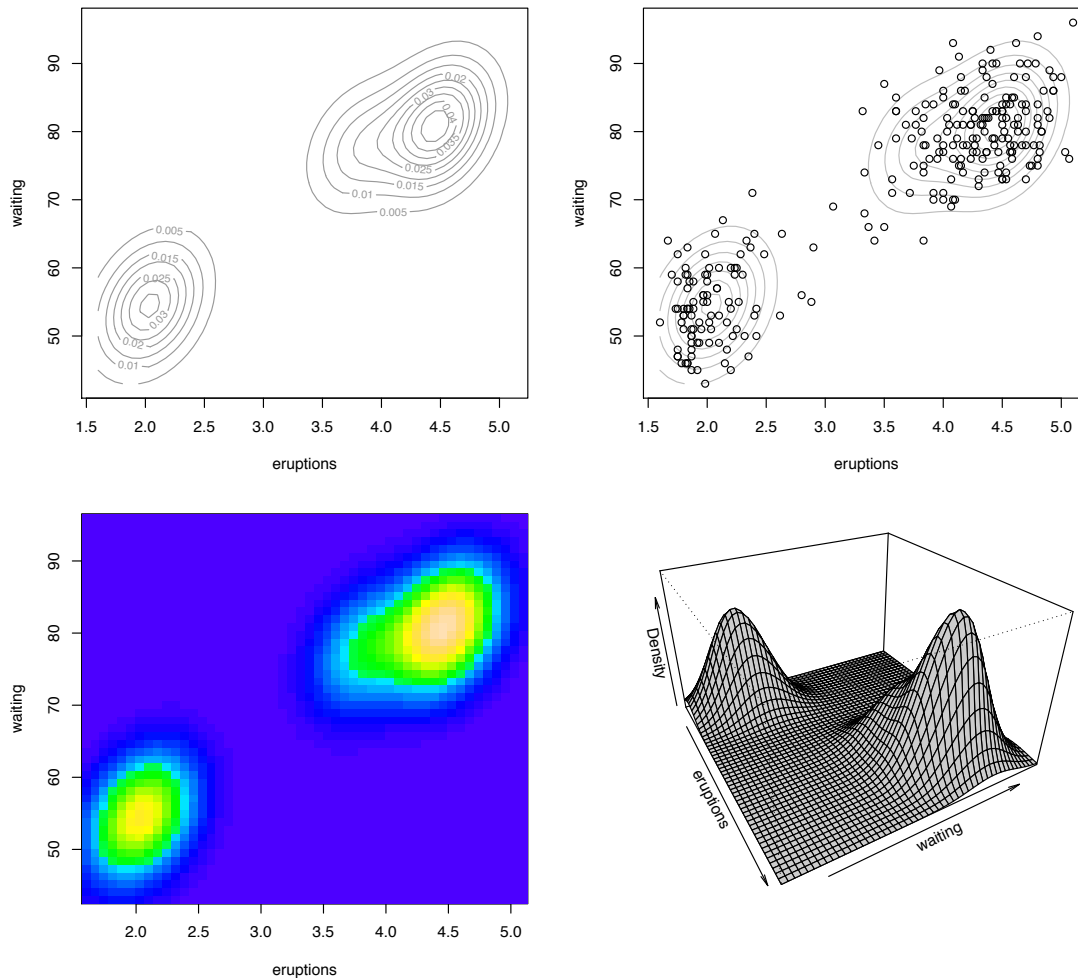
Figure 15: Contour, image and perspective plots of an `mclust` density estimate for the bivariate `faithful` dataset.

Density estimation on more than two dimensions is straightforward with `mclust`. The only difficulty is found when we want to plot the estimated density. `mclust` takes a simple route, i.e., it provides density (countour, image, and perspective) plots for all pairs of variables arranged in a matrix.

Finally, we note that the function `dens` is provided to easily get the density of a given set of points relative to an estimated model. See `help(dens)` and the examples therein.

7 Discriminant Analysis

In discriminant analysis, observations of known classification are used to classify others. `mclust` provides the function `MclustDA` which allows several approaches to discriminant analysis. We demonstrate some possible methods applied to the well-known `iris` dataset split into a training set and test set:

```
> odd <- seq(from=1, to=nrow(iris), by=2)
> trainData <- iris[odd,1:4]
```

```
> trainClass <- iris$Species[odd]
> testData <- iris[-odd,1:4]
> testClass <- iris$Species[-odd]
```

7.1 Discriminant Analysis Through Eigenvalue Decomposition

Bensmail and Celeux [4] proposed the use of Gaussian finite mixture modeling for discriminant analysis in which each known class is modeled by a single Gaussian term with the same covariance structure among classes. They named this procedure Eigenvalue Decomposition Discriminant Analysis or EDDA. This method for discriminant analysis is available as an option (named EDDA) in function MclustDA.

Consider the following model:

```
> irisEEE <- MclustDA(trainData, trainClass, modelType = "EDDA", modelNames = "EEE")
> summary(irisEEE, parameters = TRUE)
------------------------------------------------
Gaussian finite mixture model for classification
------------------------------------------------

EDDA model summary:

 log.likelihood  n df      BIC
       -125.44 75 22 -345.87

Classes        n Model G
  setosa      25   EEE 1
  versicolor 25   EEE 1
  virginica  25   EEE 1

Estimated parameters:

Class = setosa

Mixing probabilities: 1

Means:
             [,1]
Sepal.Length 5.024
Sepal.Width  3.480
Petal.Length 1.456
Petal.Width  0.228

Variances:
[,,1]
            Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length    0.264181    0.062448     0.159355    0.031413
Sepal.Width     0.062448    0.096309     0.033269    0.032224
Petal.Length    0.159355    0.033269     0.182368    0.040917
Petal.Width     0.031413    0.032224     0.040917    0.038912

Class = versicolor

Mixing probabilities: 1
```

```
Means:
              [,1]
Sepal.Length 5.992
Sepal.Width  2.776
Petal.Length 4.308
Petal.Width  1.352

Variances:
[,,1]
            Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length    0.264181    0.062448     0.159355    0.031413
Sepal.Width     0.062448    0.096309     0.033269    0.032224
Petal.Length    0.159355    0.033269     0.182368    0.040917
Petal.Width     0.031413    0.032224     0.040917    0.038912


Class = virginica

Mixing probabilities: 1

Means:
              [,1]
Sepal.Length 6.504
Sepal.Width  2.936
Petal.Length 5.564
Petal.Width  2.076

Variances:
[,,1]
            Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length    0.264181    0.062448     0.159355    0.031413
Sepal.Width     0.062448    0.096309     0.033269    0.032224
Petal.Length    0.159355    0.033269     0.182368    0.040917
Petal.Width     0.031413    0.032224     0.040917    0.038912

Training classification summary:

           Predicted
Class       setosa versicolor virginica
  setosa        25          0         0
  versicolor     0         24         1
  virginica      0          1        24

Training error = 0.026667
```

The above model fits a separate mean vector for each class, but the same ellipsoidal covariance matrix, which is essentially equivalent to linear discriminant analysis.

If we want to allow for different within-class covariance matrices (which is essentially equivalent to quadratic discriminant analysis) we can use:

```
> irisVVV <- MclustDA(trainData, trainClass, modelType = "EDDA", modelNames = "VVV")
> summary(irisVVV, parameters = TRUE)
-------------------------------------------------
Gaussian finite mixture model for classification
```

```
------------------------------------------------

EDDA model summary:

 log.likelihood  n df      BIC
       -85.149 75 42 -351.63

Classes         n Model G
  setosa      25   VVV 1
  versicolor 25   VVV 1
  virginica  25   VVV 1

Estimated parameters:

Class = setosa

Mixing probabilities: 1

Means:
              [,1]
Sepal.Length 5.024
Sepal.Width  3.480
Petal.Length 1.456
Petal.Width  0.228


Variances:
[,,1]
             Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length     0.146624     0.08848     0.014656    0.004528
Sepal.Width      0.088480     0.10160     0.005520    0.005760
Petal.Length     0.014656     0.00552     0.040864    0.003632
Petal.Width      0.004528     0.00576     0.003632    0.006016


Class = versicolor

Mixing probabilities: 1

Means:
              [,1]
Sepal.Length 5.992
Sepal.Width  2.776
Petal.Length 4.308
Petal.Width  1.352

Variances:
[,,1]
             Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length     0.296736    0.079808     0.191264    0.048816
Sepal.Width      0.079808    0.108224     0.082992    0.042848
Petal.Length     0.191264    0.082992     0.220736    0.069584
Petal.Width      0.048816    0.042848     0.069584    0.036096


Class = virginica
```

```
Mixing probabilities: 1

Means:
              [,1]
Sepal.Length 6.504
Sepal.Width  2.936
Petal.Length 5.564
Petal.Width  2.076


Variances:
[,,1]
           Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length     0.349184    0.019056     0.272144    0.040896
Sepal.Width      0.019056    0.079104     0.011296    0.048064
Petal.Length     0.272144    0.011296     0.285504    0.049536
Petal.Width      0.040896    0.048064     0.049536    0.074624

Training classification summary:

            Predicted
Class       setosa versicolor virginica
  setosa        25          0         0
  versicolor     0         24         1
  virginica      0          0        25


Training error = 0.013333
```

Other possible parameterization of the covariance matrix are reported in Table 1. If the argument `modelNames` is not provided, the best model is chosen from among all available parameterizations using the BIC. Most commonly in classification problems, model selection would be done by computing the classification error using cross-validation or, if available, a test set. mclust provides the function `cv.MclustDA` for computing k-fold cross-validation error, and the function `predict.MclustDA` for classifying new observations. With the above example data, we may compute the three statistics and compare the results as follows:

```
> models <- mclust.options()$emModelNames
> tab <- matrix(NA, nrow = length(models), ncol = 3)
> rownames(tab) <- models
> colnames(tab) <- c("BIC", "10-fold CV", "Test error")
> for(i in seq(models))
    {
      mod <- MclustDA(trainData, trainClass,
                   modelType = "EDDA", modelNames = models[i])
      tab[i,1] <- mod$bic
      tab[i,2] <- cv.MclustDA(mod, nfold = 10, verbose = FALSE)$error
      pred <- predict(mod, testData)
      tab[i,3] <- classError(pred$classification, testClass)$errorRate
    }
> tab
        BIC 10-fold CV Test error
EII -462.26   0.053333   0.066667
VII -450.76   0.053333   0.066667
EEI -417.16   0.053333   0.040000
```

```
VEI -401.31   0.040000   0.040000
EVI -424.79   0.040000   0.040000
VVI -398.00   0.040000   0.040000
EEE -345.87   0.026667   0.040000
EEV -347.45   0.026667   0.053333
VEV -331.30   0.040000   0.040000
VVV -351.63   0.026667   0.040000
```

The largest BIC is achieved by the `VEV` model, which is also one of the models with both the smallest cross-validation error and test errror. Although in this case cross-validation, test error, and BIC happen to choose the same model, for other datasets the models selected, and hence the discriminant results, could be different.

7.2 Model-based Discriminant Analysis via MclustDA

In Section 7.1, discriminant analysis was accomplished by modeling the training data with a mixture density having a single Gaussian component for each class. A more flexible alternative is to use model-based clustering to fit a Gaussian mixture model as a density estimate for each class in the training set [13].

In mclust this model can be fitted as follows:

```
> irisMCLUSTDA <- MclustDA(trainData, trainClass, modelType = "MclustDA")
```

For clarity of exposition we specified `modelType = "MclustDA"`, but this argument can be omitted since it is the default. Other arguments that could have been given are `G`, for the number of mixture components, and `modelNames`, for the parameterization of class covariance matrix. If `G` is not provided up to 5 components are fit for each class. If `modelNames` is not provided the best model for each class is chosen from among all available parameterizations using the BIC. If provided as single vector of values, for example

```
> MclustDA(trainData, trainClass, modelType = "MclustDA",
          G = 1:2, modelNames = c("EII", "EEE"))
```

the search for the largest BIC values within each class is performed among the models `EII` and `EEE`, with one or two mixture components. Such arguments could also be supplied as lists, for example

```
> MclustDA(trainData, trainClass, modelType = "MclustDA",
          G = list(1:3, 1:2, 1), modelNames = list("EII", "EEE", c("EEE", "VVV")))
```

In this case, each pair of (i) number of mixture components and (ii) model name is used in fitting the corresponding mixture model within each class. Thus, the first class is fitted with model `EII` having from one to three components, the second class with model `EEE` having from one to two components, and the last class with a single component mixture chosen between models `EEE` and `VVV`.

A fitted model can be summarized as follows:

```
> summary(irisMCLUSTDA)
------------------------------------------------
Gaussian finite mixture model for classification
------------------------------------------------

MclustDA model summary:
```

28

```
 log.likelihood  n df      BIC
       -63.55 75 53 -355.93


Classes       n Model G
  setosa     25   VEI 2
  versicolor 25   EEV 2
  virginica  25   XXX 1


Training classification summary:

           Predicted
Class         setosa versicolor virginica
  setosa          25          0         0
  versicolor       0         25         0
  virginica        0          0        25


Training error = 0
```

In this case BIC selects the `VEI` model with two components for the setosa Irises, the `EEV` model with two components for the versicolor Irises, and the single component with full covariance matrix for the virginica Irises. The `summary` method prints also the confusion table and the error for the training data. The latter turns out to be zero, but a more realistic estimate of the classification error requires the evaluation of the fitted model in a cross-validated or test set.

The function `cv.MclustDA` can be used to obtain the k-fold cross-validation classification, error rate, and standard error of classification:

```
> cv = cv.MclustDA(irisMCLUSTDA, nfold = 10, verbose = FALSE)
> str(cv)
List of 3
 $ classification: Factor w/ 3 levels "setosa","versicolor",..: 1 1 1 1 1 1 1 1 1 1 ...
 $ error         : num 0.04
 $ se            : num 0.0194
```

Given a test data set and the corresponding classes for the observations in the test set, the test error rate is obtained as follows:

```
> pred <- predict(irisMCLUSTDA, testData)
> classError(pred$classification, testClass)$errorRate
[1] 0.013333
```

Alternatively, the `summary` method could be used to obtain the test error rate:

```
> summary(irisMCLUSTDA, newdata = testData, newclass = testClass)
------------------------------------------------
Gaussian finite mixture model for classification
------------------------------------------------

MclustDA model summary:

 log.likelihood  n df      BIC
       -63.55 75 53 -355.93


Classes       n Model G
```

```
  setosa     25   VEI 2
  versicolor 25   EEV 2
  virginica  25   XXX 1
```

```
Training classification summary:
```

```
          Predicted
Class          setosa versicolor virginica
  setosa           25          0         0
  versicolor        0         25         0
  virginica         0          0        25
```

```
Training error = 0
```

```
Test classification summary:
```

```
          Predicted
Class          setosa versicolor virginica
  setosa           25          0         0
  versicolor        0         24         1
  virginica         0          0        25
```

```
Test error = 0.013333
```

7.3   Plotting methods

Several plots for discriminant analysis objects are available throught the `plot` method. Here we illustrate a subset of them, inviting the interested reader to consult `help(plot.MclustDA)` for a more comprehensive list and examples.

The default plot, corresponding to the omitted argument `what = "scatterplot"`, is obtained as:

```
> plot(irisMCLUSTDA)
```

and it is shown in Figure 16. This produces a scatterplot matrix for pairs of variables with points marked according to the corresponding class, and mixture components ellipses superimposed. A subset of variables can be selected with the argument `dimens`, for instance:

```
> plot(irisMCLUSTDA, dimens = 3:4)
```

A plot of data points, either from train set and from test set, marked according the known or estimated classification is easily obtained as:

```
> plot(irisMCLUSTDA, what = "classification", dimens = 3:4)
> plot(irisMCLUSTDA, what = "classification", newdata = testData, dimens = 3:4)
```

where we selected only the third and fourth variable to plot. The resulting graphs are shown in top panels of Figure 17.

Two other available displays are:

```
> plot(irisMCLUSTDA, what = "train&test", newdata = testData, dimens = 3:4)
> plot(irisMCLUSTDA, what = "error", newdata = testData, newclass = testClass, dimens = 3:4)
```
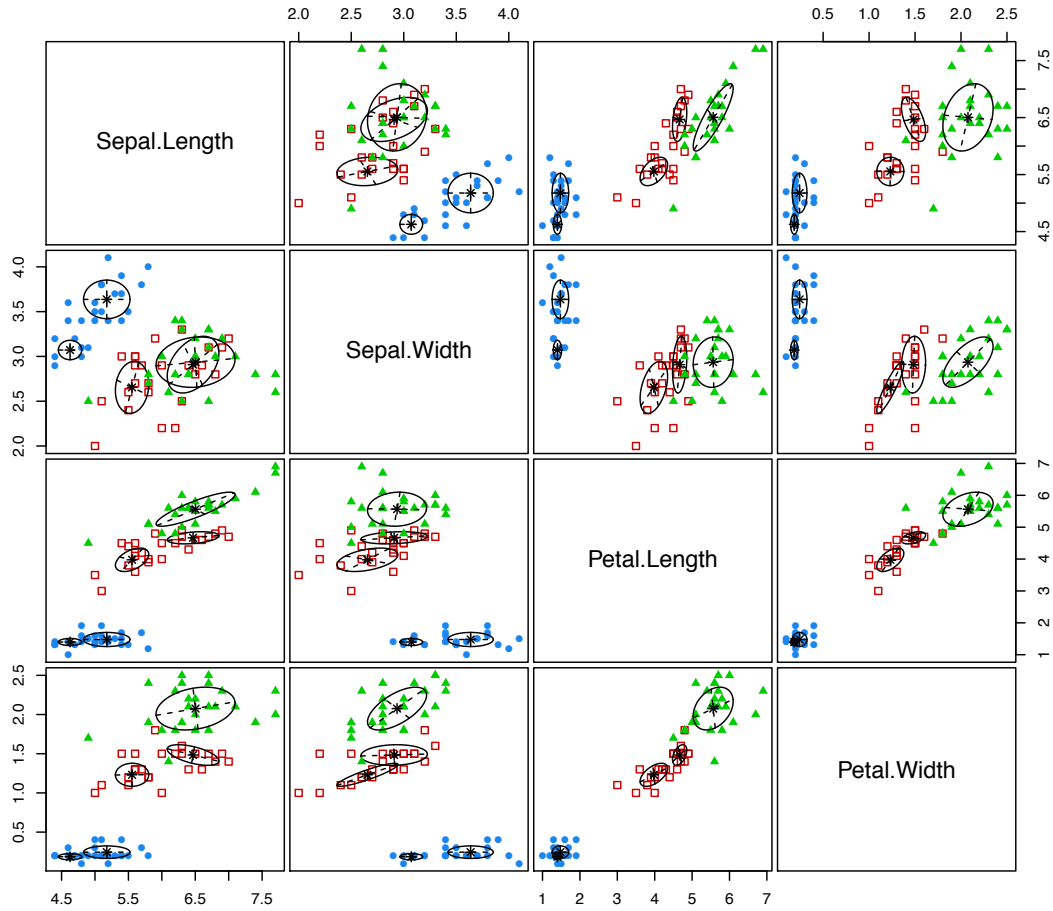
Figure 16: Scatterplot matrix for variables with points marked according to the corresponding class, and with mixture components ellipses from `MclustDA` superimposed for the `iris` dataset.

In the first graph we plot both the training data and the test set with different plotting symbols and colors, while in the second case the test set is plotted with the misclassified observation identified by a black square. These graphs are shown in the bottom panels of Figure 17.

```
> par(mfrow = c(2,2), mar = c(5,4,2,1))
> plot(irisMCLUSTDA, what = "classification", dimens = 3:4)
> plot(irisMCLUSTDA, what = "classification", newdata = testData, dimens = 3:4)
> plot(irisMCLUSTDA, what = "train&test", newdata = testData, dimens = 3:4)
> plot(irisMCLUSTDA, what = "error", newdata = testData, newclass = testClass, dimens = 3:4)
```

7.4   Classification of univariate data

To illustrate discriminant analysis on univariate data, we use simulated data from a normal mixture consisting of two components with variance 1 centered at -9 and 9, respectively, and one component with variance 4 centered at 0:
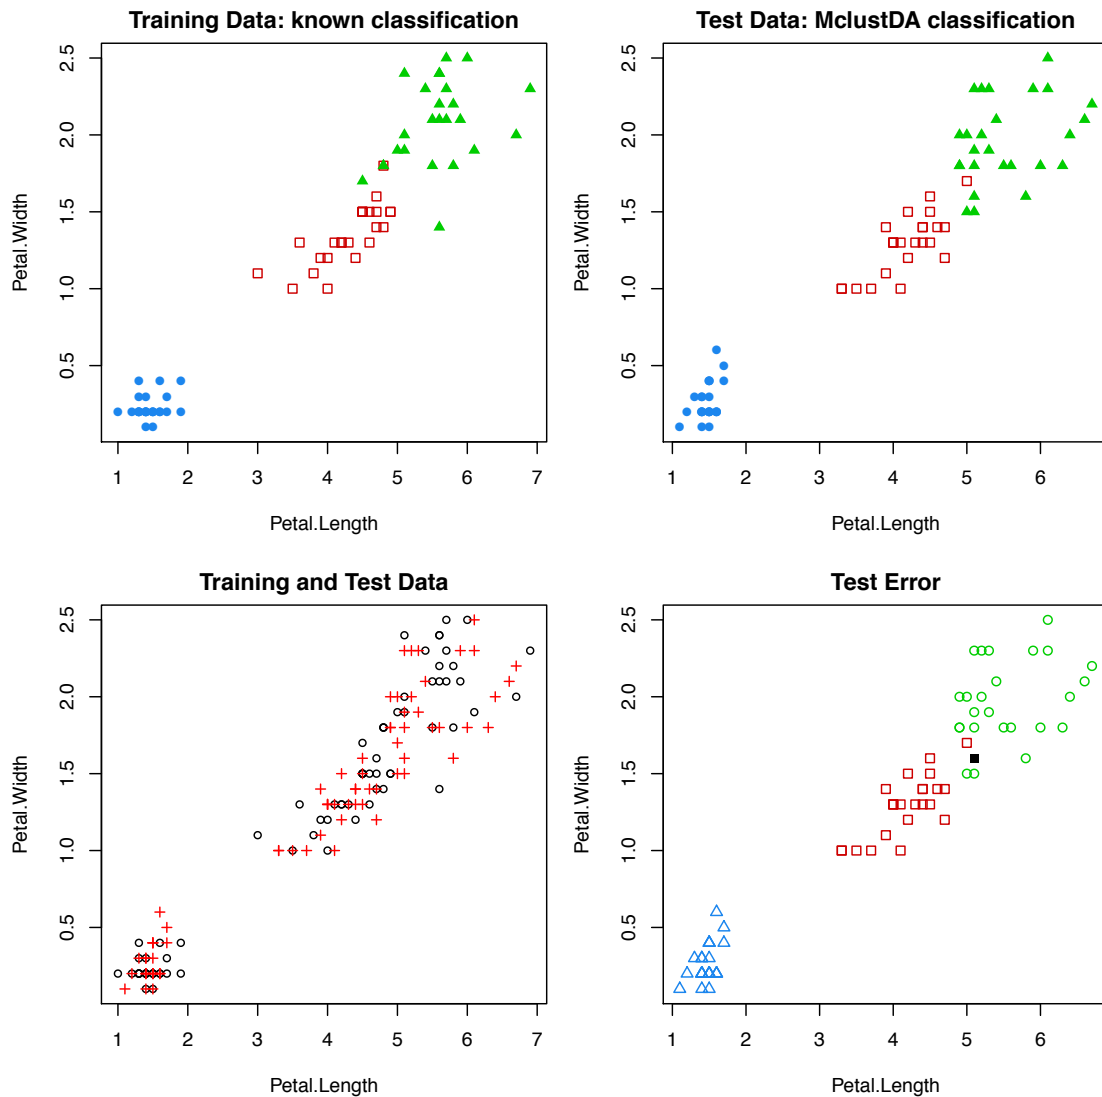
Figure 17: Classification plots for a selected pair of variables with points marked according to the known or estimated classification from `MclustDA` for the `iris` dataset.

```
> set.seed(0)
> x <- c(rnorm(300, -9), rnorm(400, 0, sd = 2), rnorm(300, 9))
> xClass <- rep(1:3, times = c(300,400,300))
```

We use the following simulated data as a test set:

```
> set.seed(1)
> y <- c(rnorm(100, -9), rnorm(100, 0, sd = 2), rnorm(100, 9))
> yClass <- rep(1:3, each = 100)
```

In both cases we set the seed of the random number generator so the exact example can be reproduced.

In discriminant analysis with EDDA (Section 7.1) we assume that each component constitutes a separate group. We may use either leave-one-out crossvalidation or BIC for choosing between the only two possible models for univariate data:

```
> cv1EMtrain(x,labels=xClass)
    E     V
0.006 0.002

> bicEMtrain(x,labels=xClass)
      E        V
-5786.7 -5607.2
```

The varying variance model V was selected as the optimal model for the training set under both criteria. The training and test errors for the data with this model are as follows:

```
> modV <- MclustDA(x, xClass, modelType = "EDDA", modelName = "V")
> summary(modV, newdata = y, newclass = yClass)
------------------------------------------------
Gaussian finite mixture model for classification
------------------------------------------------

EDDA model summary:

 log.likelihood    n df      BIC
          -2776 1000  6 -5593.4

Classes   n Model G
      1 300     V 1
      2 400     V 1
      3 300     V 1

Training classification summary:

      Predicted
Class   1   2   3
    1 300   0   0
    2   0 399   1
    3   0   1 299

Training error = 0.002

Test classification summary:

      Predicted
Class   1   2   3
    1 100   0   0
    2   0 100   0
    3   0   0 100

Test error = 0
```

The classification and classification errors for the training data are shown in Figure 18 and have been obtained with the following code:

```
> par(mfrow = c(1,3), mar = c(4,4,2,1))
> plot(modV)
> plot(modV, what = "classification")
> plot(modV, what = "error")
```
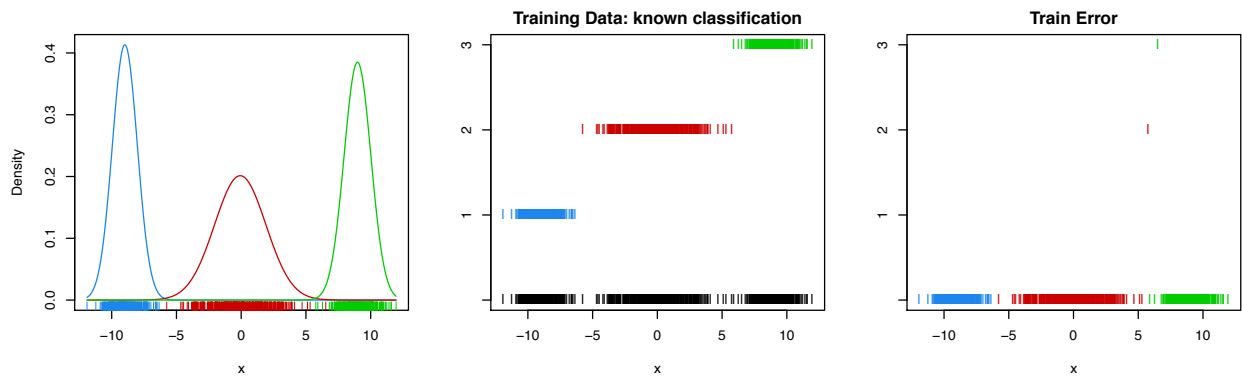
Figure 18: LEFT: Densities for the univariate simulation data. There are two components with variance 1 centered at -9 and 9, respectively, and one component with variance 4 centered at 0. CENTER: Training step classification. RIGHT: Misclassified training observations.

To illustrate the discriminant analysis via the MclustDA methodology (see Section 7.2) we use the same simulated univariate data as above but assume that observations are grouped by component variance:

```
> xClass <- rep(c(1,2,1), times = c(300,400,300))
> yClass <- rep(c(1,2,1), each = 100)
```

The training stage fits a two component equal-variance model to one group, and a one-component model to the other:

```
> modMCLUSTDA <- MclustDA(x, xClass)
> summary(modMCLUSTDA)
-------------------------------------------------
Gaussian finite mixture model for classification
-------------------------------------------------

MclustDA model summary:

 log.likelihood    n df      BIC
        -2863.5 1000  6 -5768.5


Classes    n Model G
      1 600     E 2
      2 400     X 1

Training classification summary:

     Predicted
Class   1   2
    1 599   1
    2   1 399


Training error = 0.002
```

The classification error rates are the same as we obtained for discriminant analysis via EM with the 3-class grouping.

8 Displays for Multidimensional Data

Once parameter values of a mixture model fit are available, projections of the data showing the means and standard deviations of the corresponding components or clusters may be plotted. For bivariate data, density and uncertainty surfaces may also be plotted.

8.1 Displays for Bivariate Data

The function `mclust2Dplot` may be used for displaying the classification, uncertainty or classification errors for `mclust` models of bivariate data. In the following example, classification and uncertainty plots are produced for the `faithful` dataset in Figure 1.

```
> faithfulMclust <- Mclust(faithful)
```

```
> mclust2Dplot(data = faithful, what = "classification", identify = TRUE,
               parameters = faithfulMclust$parameters, z = faithfulMclust$z)
```

```
> mclust2Dplot(data = faithful, what = "uncertainty", identify = TRUE,
               parameters = faithfulMclust$parameters, z = faithfulMclust$z)
```

The resulting plots are displayed in Figure 19.
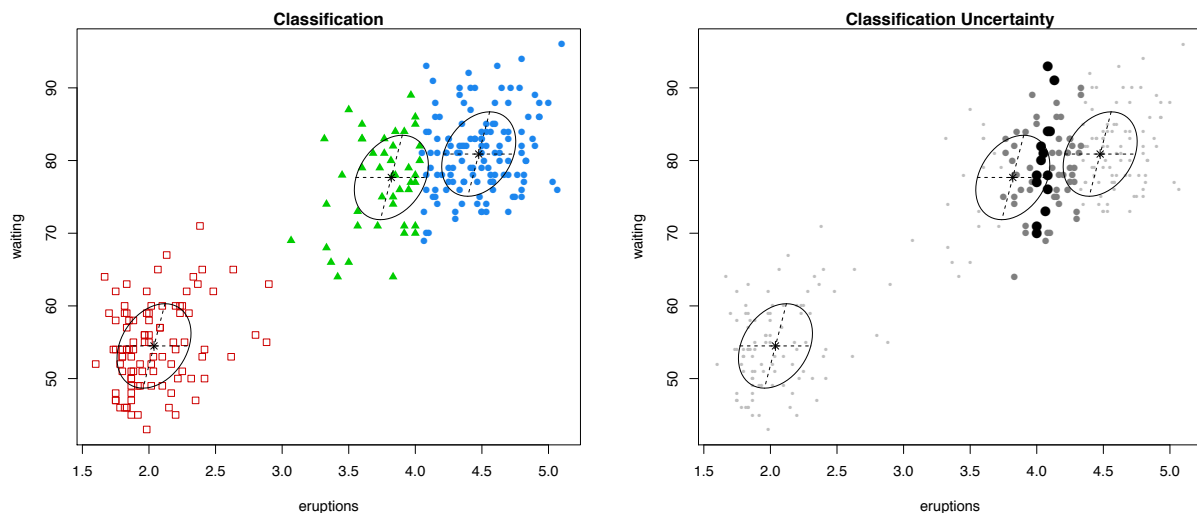
Figure 19: Classification (left) and uncertainty (right) plots created with `mclust2Dplot` for the `Mclust` model of the `faithful` dataset. The ellipses shown are the multivariate analogs of the standard deviations for each mixture component. In the classification plot, points in different classes are indicated by different symbols. In the uncertainty plot, the symbols have the following meaning: large filled symbols, 95% quantile of uncertainty; smaller open symbols, 75–95% quantile; small dots, first three quartiles of uncertainty.

The function `surfacePlot` may be used for displaying the density or uncertainty for `mclust` models of bivariate data. It also returns the grid coordinates and corresponding surface values. The following example shows how to display density and uncertainty surfaces for the `Mclust` model fit to the `faithful` dataset.

```
> surfacePlot(data = faithful, what = "density", type = "contour",
              parameters = faithfulMclust$parameters, transformation = "sqrt")
```

```
> surfacePlot(data = faithful, what = "uncertainty", type = "image",
              parameters = faithfulMclust$parameters,  transformation = "log")
```

The resulting plots are displayed in Figure 20.

Figure 20: Density (left column) and uncertainty (right column) surfaces for the `faithful` dataset. A square root transformation was used for the density plot, which is plotted as a contour surface. A logarithmic transformation was used for the uncertainty plot, which is plotted as an image surface.

8.2 Displays for Higher Dimensional Data

8.2.1 Coordinate Projections

To plot coordinate projections in `mclust` use the function `coordProj`. The example we consider is a 3-group model for the `iris` dataset:

```
> irisBIC <- mclustBIC(iris[,-5])
> irisSummary3 <- summary(irisBIC, data = iris[,-5], G = 3)

> coordProj(data = iris[,-5], dimens = c(2,4), what = "classification",
            parameters = irisSummary3$parameters, z = irisSummary3$z)
> coordProj(data = iris[,-5], dimens = c(2,4), what = "uncertainty",
            parameters = irisSummary3$parameters, z = irisSummary3$z)
> coordProj(data = iris[,-5], dimens = c(2,4), what = "errors",
            parameters = irisSummary3$parameters, z = irisSummary3$z, truth = iris[,5])
```
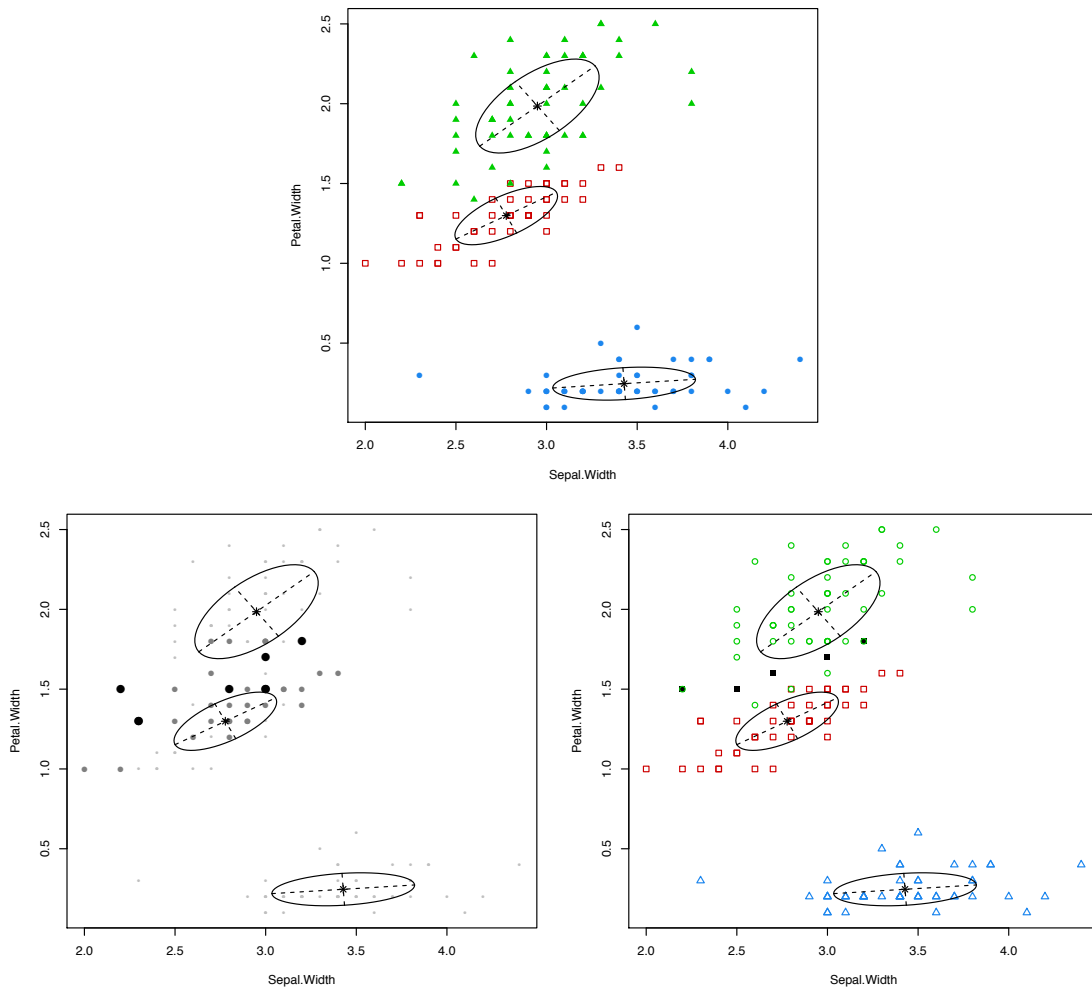

Figure 21: A coordinate projection of the `iris` dataset created with `coordProj`. Plots show the 3-group model-based classification (top) with associated uncertainty (bottom, left) and classification errors (bottom,right).

These plots are displayed in Figure 21.

37

8.2.2 Random Projections

To plot random projections in mclust use the function randProj. Again we consider is a 3-group model for the iris dataset:

```
> randProj(data = iris[,-5], seed = 43, what = "classification",
          parameters = irisSummary3$parameters, z = irisSummary3$z)
> randProj(data = iris[,-5], seed = 79, what = "classification",
          parameters = irisSummary3$parameters, z = irisSummary3$z)
> randProj(data = iris[,-5], seed = 201, what = "classification",
          parameters = irisSummary3$parameters, z = irisSummary3$z)
```

These plots are displayed in Figure 22.

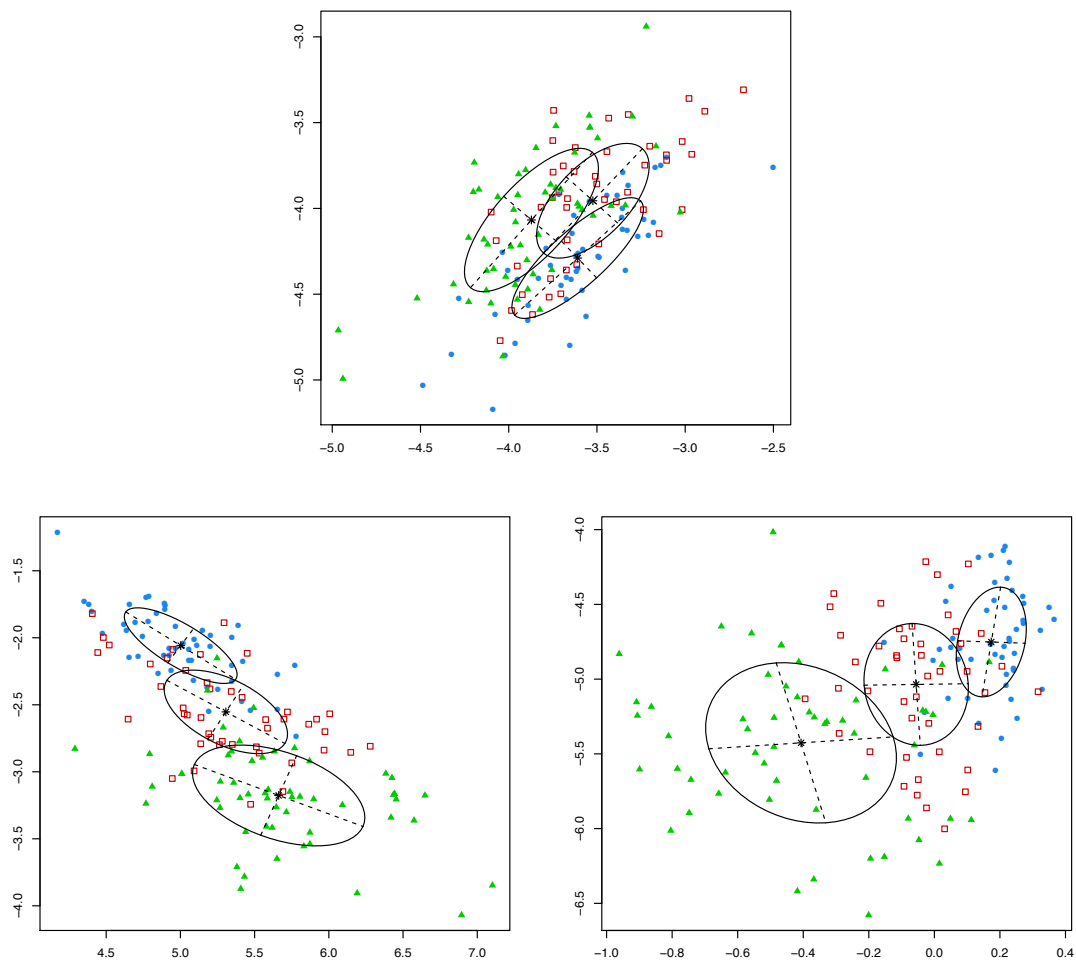**Figure 22:** Some random projections of the iris dataset created with randProj. Plots show the 3-group classification from model-based clustering with three different seeds.

8.3 Dimension reduction for model-based clustering and classification

New functionality has been added to mclust package for data projection on a dimension reduced subspace [31]. This dimension reduction method allows the visualization of the clustering or classification structure obtained from a Gaussian finite mixture model.

8.3.1 Clustering

Consider the following clustering scenario:

```
> MCLUST = Mclust(iris[,1:4])
> MCLUST
'Mclust' model object:
 best model: ellipsoidal, equal shape (VEV) with 2 components
```

Once a mixture model has been fitted, we may want to look at a plot in a reduced subspace which capture most of the clustering structure contained in the data. The function MclustDR allows to estimate the basis of such subspace:

```
> DR = MclustDR(MCLUST)
> summary(DR)
-----------------------------------------------------------------
Dimension reduction for model-based clustering and classification
-----------------------------------------------------------------

Mixture model type: Mclust

Estimated basis vectors:
                 Dir1     Dir2      Dir3      Dir4
Sepal.Length  0.18224 -0.23497   0.53649  0.030665
Sepal.Width   0.73764  0.16437  -0.36366  0.902191
Petal.Length -0.60349 -0.27048  -0.49512  0.274447
Petal.Width  -0.24183  0.91903   0.57861 -0.331346


                 Dir1     Dir2      Dir3      Dir4
Eigenvalues   0.85068  0.34586  0.064915   0.030269
Cum. %       65.85600 92.63125 97.656678 100.000000
```

The directions which span the reduced subspace are defined as a set of linear combinations of the original features, ordered by importance as quantified by the associated eigenvalues. In the example, the first two directions account for most of the clustering structure as shown in Figure 23. The graph is obtained with the following code:

```
> plot(DR, what = "evalues")
```

The main contribution to the first direction is from differences in cluster means, while for the second direction is from differences in cluster variances.

A scatterplot matrix of data projected onto the estimated subspace is shown in Figure 24 and it is obtained as follows:

```
> plot(DR)
```

Figure 25 contains examples of bivariate plots. The left graph is a bivariate contour plot of mixture components densities:
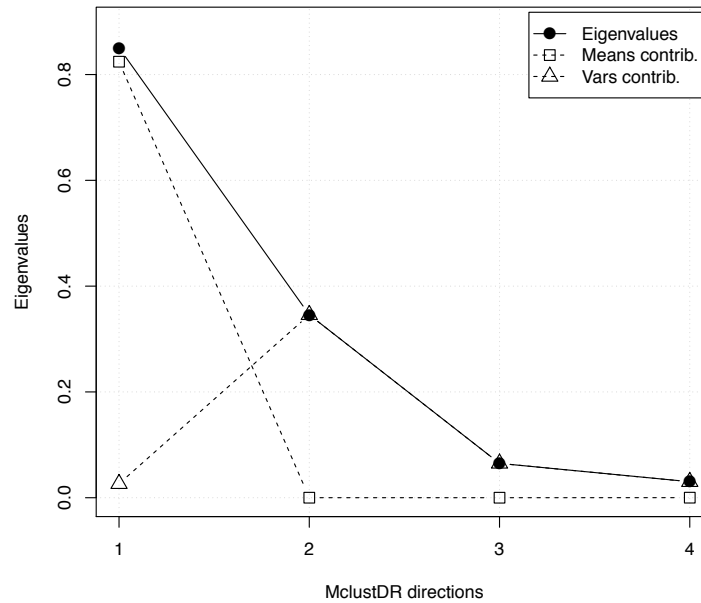
Figure 23: Plot of eigenvalues from `MclustDR` applied to a clustering model for the `iris` dataset.

```
> plot(DR, what = "contour")
```

while the graph on the right is a scatterplot with uncertainty boundaries:

```
> plot(DR, what = "boundaries")
```

where the uncertainty is shown using a greyscale with darker regions indicating higher uncertainty. Both plot projects the data onto the first two directions, i.e., using the default argument `dimens = c(1,2)`. Setting this argument to any other pair of values produces a projection along different directions.

Marginal distributions along selected directions can be obtained as follows:

```
> plot(DR, what = "density", dimens = 1)
> plot(DR, what = "density", dimens = 2)
```

and they are shown in Figure 26.

Many other graphs can be obtained by specifying the argument `what`. The interested reader may consult `help(plot.MclustDR)` for a comprehensive list and examples.

8.3.2 Classification

The method illustrated for clustering can also be applied to a classification scenario, except that now we must take into account that known classes can be made of one or more mixture components. For example, consider the `banknote` dataset contained in `mclust`:

```
> data(banknote)
> MCLUSTDA = MclustDA(data = banknote[,-1], class = banknote$Status)
> summary(MCLUSTDA)
--------------------------------------------------
Gaussian finite mixture model for classification
```
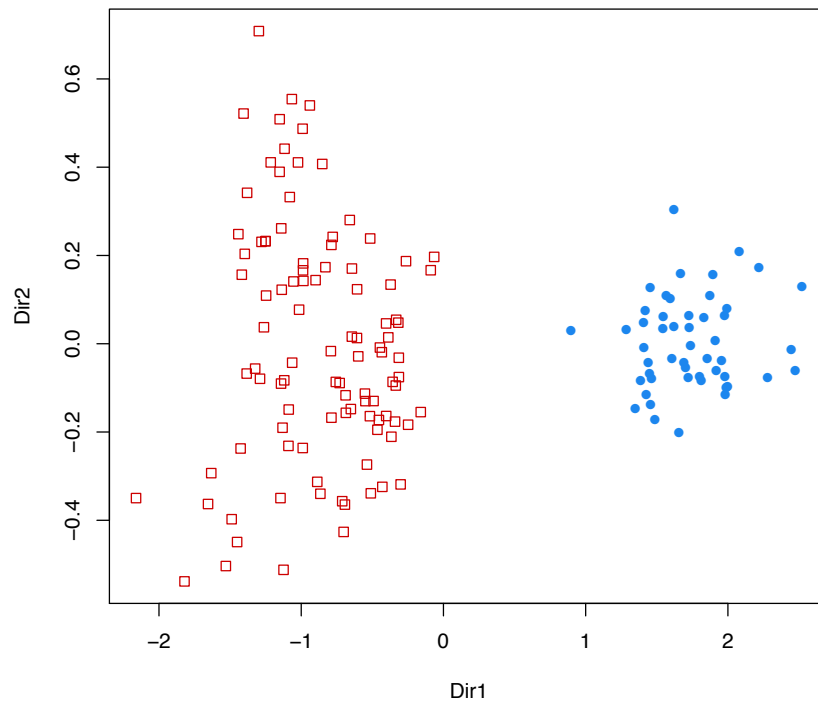
Figure 24: Scatterplot of clustering structure from `MclustDR` for the `iris` dataset.

```
--------------------------------------------------

MclustDA model summary:

 log.likelihood   n df      BIC
       -652.97 200 68 -1666.2

Classes         n Model G
  counterfeit 100   EEE 3
  genuine     100   XXX 1

Training classification summary:

              Predicted
Class         counterfeit genuine
  counterfeit         100       0
  genuine               0     100

Training error = 0

> DR = MclustDR(MCLUSTDA)
> summary(DR)
-----------------------------------------------------------------
Dimension reduction for model-based clustering and classification
-----------------------------------------------------------------

Mixture model type: MclustDA
```
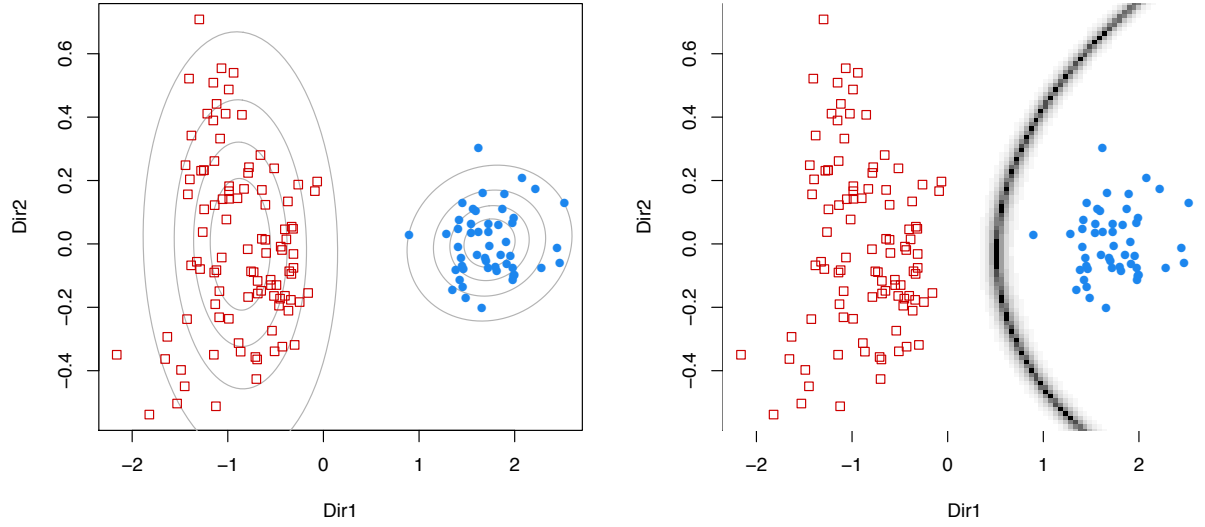
Figure 25: Contour plot of mixture densities (left) and uncertainty boundaries (right) from `MclustDR` for the `iris` dataset.
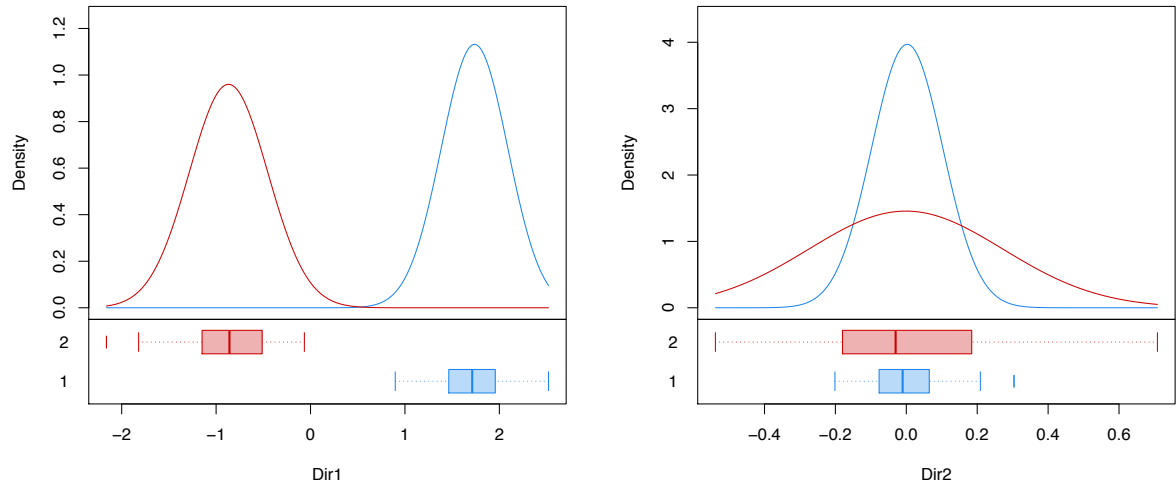
Figure 26: Marginal mixture densities for the first two directions from `MclustDR` for the `iris` dataset.

```
Estimated basis vectors:
              Dir1      Dir2       Dir3       Dir4       Dir5       Dir6
Length   -0.094293  -0.32902  -0.0043289  -0.385202   0.703088   0.551233
Left     -0.318248  -0.40252  -0.8638528   0.700590   0.492784  -0.036375
Right     0.223543  -0.12154   0.2911526  -0.587057   0.018105  -0.811847
Bottom    0.621442   0.47022  -0.0486602  -0.010436   0.111424   0.093691
Top       0.553684   0.28743  -0.2984504  -0.116710  -0.496313   0.153537
Diagonal -0.383536   0.64127  -0.2784408  -0.049267  -0.061325  -0.058068


               Dir1      Dir2      Dir3       Dir4       Dir5       Dir6
Eigenvalues  0.87532   0.60038   0.17223   0.066972    0.02501  1.8745e-03
Cum. %      50.25411  84.72339  94.61147  98.456517  99.89238  1.0000e+02
```

A pairs plot of the data projected along the first three directions is obtained with code:

```
> plot(DR, what = "pairs", dim = 1:3)
```


Figure 27: Pairs plot for the first three dimensions from `MclustDR` for the `banknote` dataset.

A bivariate contour plot of mixture components densities for each class and the classification boundaries are given by:

```
> plot(DR, what = "contour")
> plot(DR, what = "boundaries")
```

9 Combining Gaussian Mixture Components for Clustering

In traditional model-based clustering, each cluster corresponds to a single mixture component. However, it may be desirable in some situations to model a single cluster using several mixture components. The function `clustCombi` provides a means for obtaining models with clusters represented by multiple mixture components.

Consider the following simulated dataset from a bivariate Gaussian mixture distribution:

```
> data(Baudry_etal_2010_JCGS_examples)
> MCLUST <- Mclust(ex4.1)
> MCLUST
```
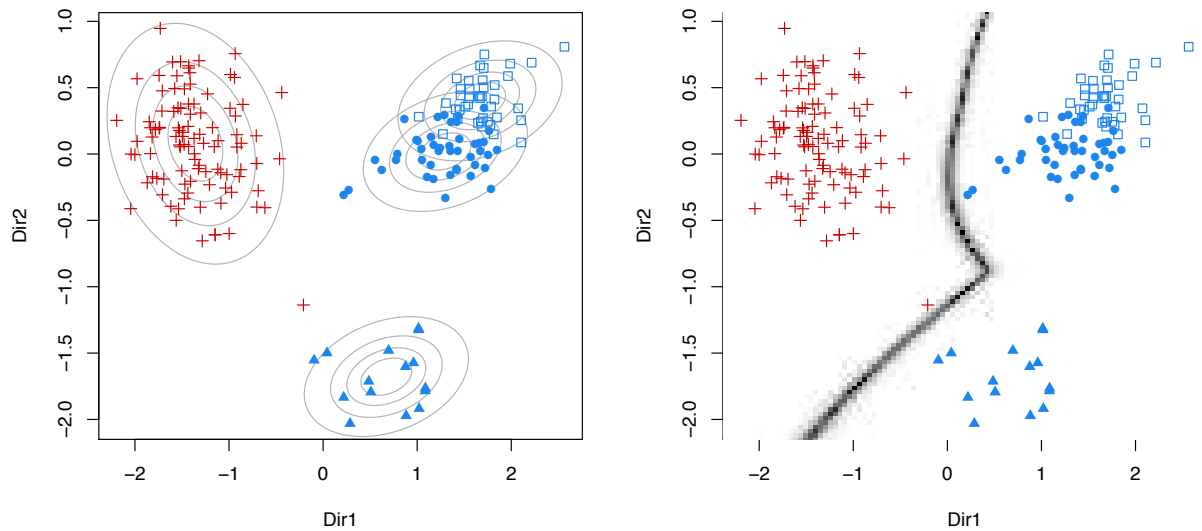
Figure 28: Contour plot of mixture densities for each class (left) and classification boundaries (right) from `MclustDR` for the `banknote` dataset.

```
'Mclust' model object:
 best model: ellipsoidal, equal volume and shape (EEV) with 6 components
```

The function `clustCombi` combines the mixture components hierarchically according to an entropy criterion, following the methodology in [3]. The solutions with numbers of classes between the number of components selected by BIC and one are returned:

```
> CLUSTCOMBI <- clustCombi(ex4.1, MCLUST)
> CLUSTCOMBI
 EM/BIC Solution
 ---------------

Number of components: 6
Model name: EEV

Component num.1:
              proportion: 0.12
              mean: 7.92  -0.02
Component num.2:
              proportion: 0.20
              mean: 8.07  4.98
Component num.3:
              proportion: 0.19
              mean: 1.01  4.96
Component num.4:
              proportion: 0.19
              mean: 1.11  5.12
Component num.5:
              proportion: 0.22
              mean: -0.02  0.06
Component num.6:
              proportion: 0.08
```

```
             mean: 8.10  0.17

Combining steps
---------------

 Step | Classes combined at this step | Classes labels after this step
------|-------------------------------|-------------------------------
   0  |              ---              | 1   2   3   4   5   6
   1  |            3  &  4            | 1   2   3   5   6
   2  |            1  &  6            | 1   2   3   5
   3  |            3  &  5            | 1   2   3
   4  |            1  &  2            | 1   3
   5  |            1  &  3            | 1

Classification for K classes: output$classification[[K]]
Combining matrix (K classes -> (K-1) classes): output$combiM[[K]]
```

The output shows the the components proportions and means of the fitted Gaussian mixture model, followed by the informations about the combining steps.
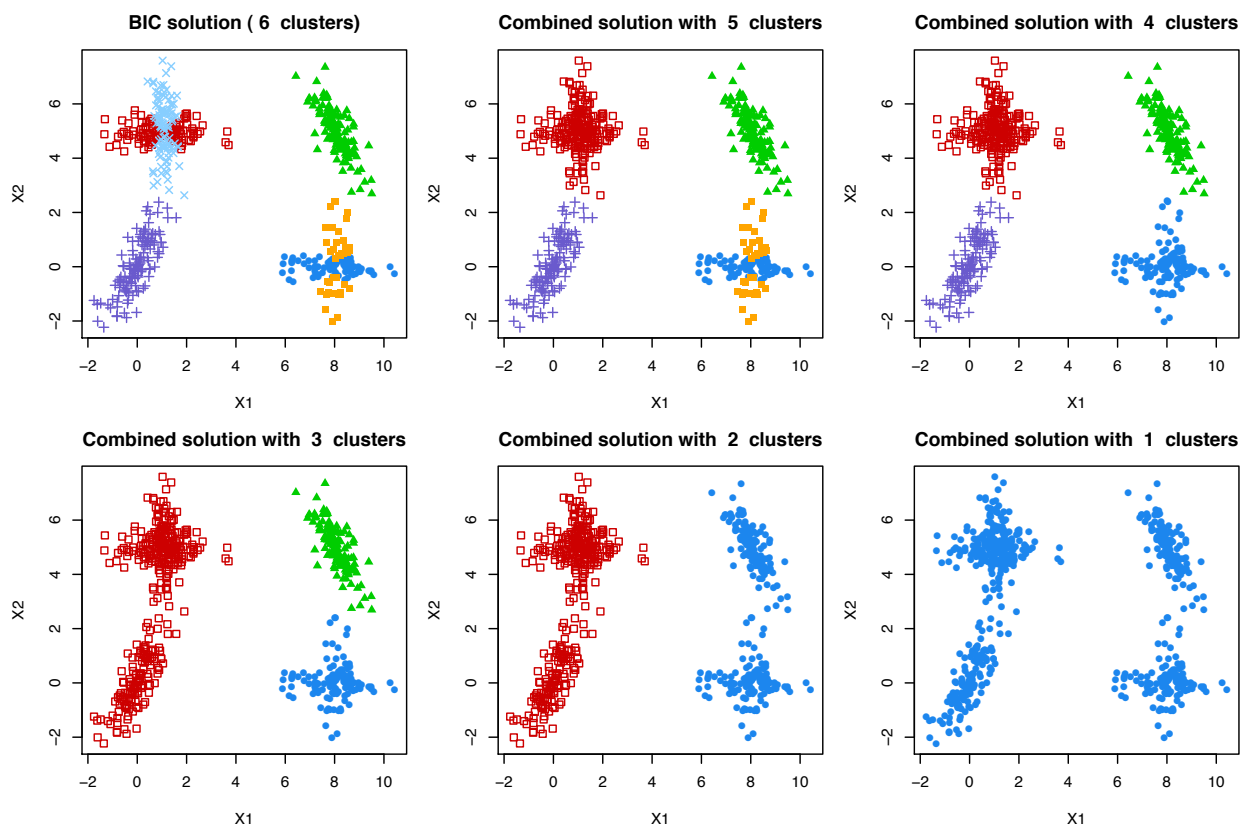
Figure 29: Solutions from combining mixture components for the `ex4.1` dataset.

The `plot` method displays the hierarchy of combined solutions and then some "entropy plots" which may help a user to select the optimal number of clusters:

```
> plot(CLUSTCOMBI, ex4.1)
```

The graphs are shown in Figure 29 and 30, with the entropy plots which suggests a four-clusters solution.
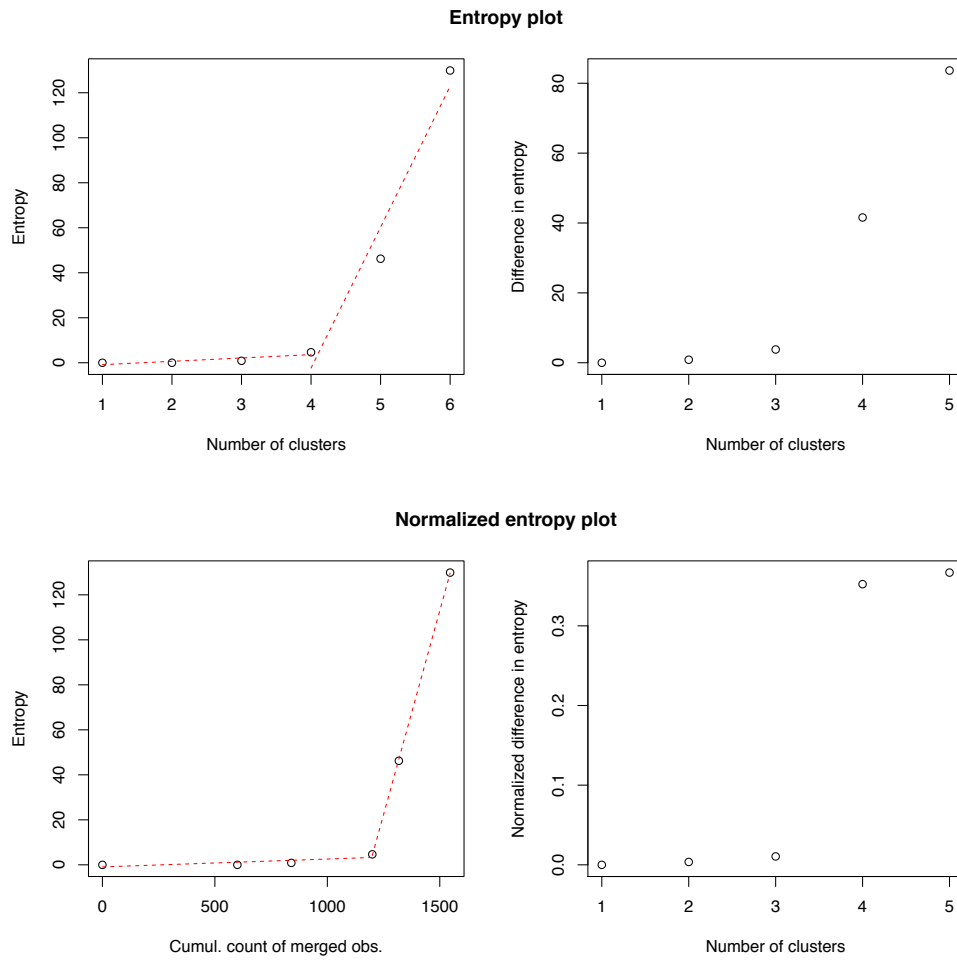
Figure 30: Entropy plots for combining mixture components for the `ex4.1` dataset.

10 Simulation from Mixture Densities

Given the parameters for a mixture model, data can be simulated from that model for evaluation and verification. The function `sim` allows simulation from mixture models generated by `mclust` functions. Besides the model, `sim` allows a seed as input for reproducibility. As an example, below we simulate two different datasets of the same size as the `faithful` dataset from the model produced by `Mclust` for the `faithful` dataset:

```
> faithfulMclust <- Mclust(faithful)
> sim0 <- sim( modelName = faithfulMclust$modelName,
               parameters = faithfulMclust$parameters,
               n = nrow(faithful), seed = 0)
> sim1 <- sim(modelName = faithfulMclust$modelName,
              parameters = faithfulMclust$parameters,
              n = nrow(faithful), seed = 1)
```

The results can be plotted as follows:

```
> xlim <- range(c(faithful[,1],sim0[,2],sim1[,2]))
> ylim <- range(c(faithful[,2],sim0[,3],sim1[,3]))
> mclust2Dplot(data=faithful, parameters = faithfulMclust$parameters,
               classification = faithfulMclust$classification, xlim = xlim, ylim = ylim)
> mclust2Dplot(data=sim0[,-1], parameters = faithfulMclust$parameters,
               classification = sim0[,1], xlim = xlim, ylim = ylim)
> mclust2Dplot(data=sim1[,-1], parameters = faithfulMclust$parameters,
               classification = sim1[,1], xlim = xlim, ylim = ylim)
```

The plots are shown in Figure 31. Note that `sim` produces a dataset in which the first column is the classification.

11 Extensions

11.1 Large Datasets

`Mclust` and `mclustBIC` include a provision for using a subsample of the data in the hierarchical clustering phase before applying EM to the full data set, in order to extend the method to larger datasets. Some other methods for handling such cases are discussed in [34, 19]. The following example uses a random sample of size 100 in the initial hierarchical clustering phase of `EMclust` applied to the iris data:

```
> nrow(iris)
[1] 150
> S <- sample(1:nrow(iris), size = 100)
> Mclust(iris[,-5], initialization = list(subset = S))
'Mclust' model object:
 best model: ellipsoidal, equal shape (VEV) with 3 components
```

For very large data sets, the capabilities of `mclust` can be easily used for classification. First, a discriminant analysis with the methodology implemented in `MclustDA` can be performed on a subset of the data. Then the remaining data points can then be classified (in reasonable sized blocks) using the `predict` method as described in section 7.
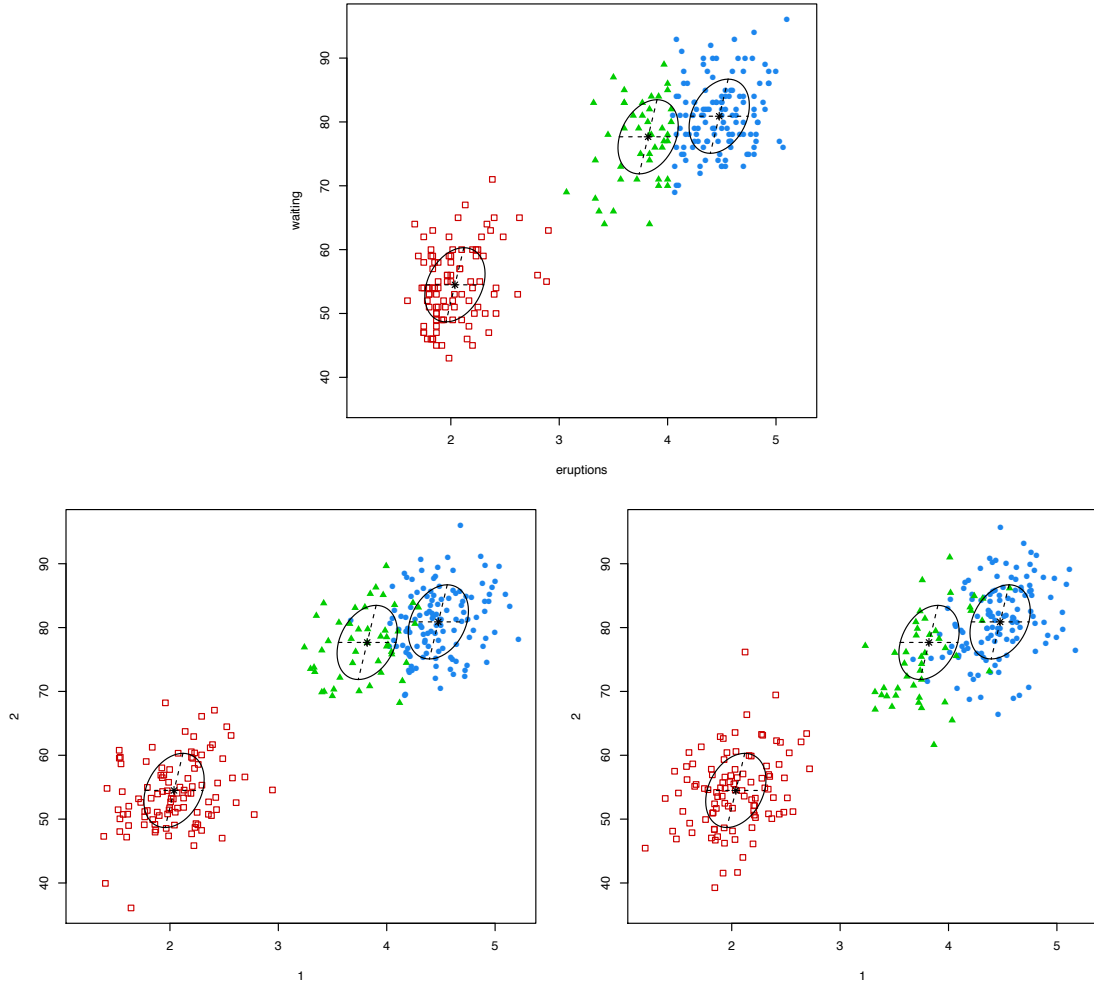
Figure 31: Data simulated from a model of the `faithful` dataset. The top hand figure is the `faithful` dataset, and the bottom figures are datasets of the same size simulated from the `Mclust` model for the `faithful` dataset. The ellipse defined by the covariance matrices of the model is shown on all of the plots.

11.2 High-Dimensional Data

Models in which the orientation is allowed to vary between clusters (`EEV`, `VEV`, `EVV`, `VVV`), have $\mathcal{O}(d^2)$ parameters per cluster, where d is the dimension of the data. For this reason, `mclust` may not work well or may otherwise be inefficient for these models when applied to high-dimensional data. It may still be possible to analyze such data with `mclust` by restriction to models with fewer parameters (e.g. spherical or diagonal models), or else by applying a dimension-reduction technique such as principal components.

Some of the more parsimonious models (e.g. spherical, diagonal, or fixed covariance) can be applied to datasets in which the number of observations is smaller than the data dimension.

11.3 Missing Data

At present, `mclust` has no direct provision for handling missing values in data. However, a function `imputeData` has been added to the `mclust` package for creating datasets with missing data imputations using the `mix` package. Here we illustrate the use of the `imputeData` to fill in missing values

in the continuous portion of the `stlouis` dataset provided with the `mix` package (we remove the first 3 categorical variables, since `mclust` is intended for continuous variables).

```
> library(mix)
> data(stlouis)
> dim(stlouis)
[1] 69  7

> head(stlouis)
     G D1 D2  R1  V1  R2  V2
4001 1 NA  1 110  NA  NA 150
4004 1  1  2 118 165  NA 130
4005 1  2 NA 116 145 114 125
4006 1 NA NA  NA  NA 126  NA
4008 1  1 NA 118 140 118 123
4009 1 NA NA  NA 120 105 138

> stlimp <- imputeData( stlouis[,-(1:3)])
Steps of EM:
1...2...3...4...5...6...7...8...9...10...11...12...13...14...15...16...17...18...19...20...21...22...2
Steps of Data Augmentation:
1...2...3...4...5...6...7...8...9...10...11...12...13...14...15...16...17...18...19...20...21...22...2
```

Note that the values obtained for the missing entries will vary depending on the random number seed set in function `imputeData`, chosen randomly by default. It is usually desirable to combine multiple imputations in analyses involving missing data. See [23] for details and references on multiple imputation.

Another function `imputePairs` has been added to the `mclust` package for visualizing missing data imputations:

```
> imputePairs(stlouis[,-(1:3)], stlimp)
```

A pairs plot showing the imputed values is displayed in Figure 32.

12 Function Summary

12.1 Hierarchical Clustering

hc	Merge sequences for model-based hierarchical clustering.
hclass	Classifications corresponding to `hc` results.

12.2 Parameterized Gaussian Mixture Models

em	EM algorithm (starting with E-step).
me	EM algorithm (starting with M-step).
me.weighted	EM algorithm with weights (starting with M-step)
estep	E-step of the EM algorithm.
mstep	M-step of the EM algorithm.
mvn	One-component fit.

12.3 Density Computation for Parameterized Gaussian Mixtures

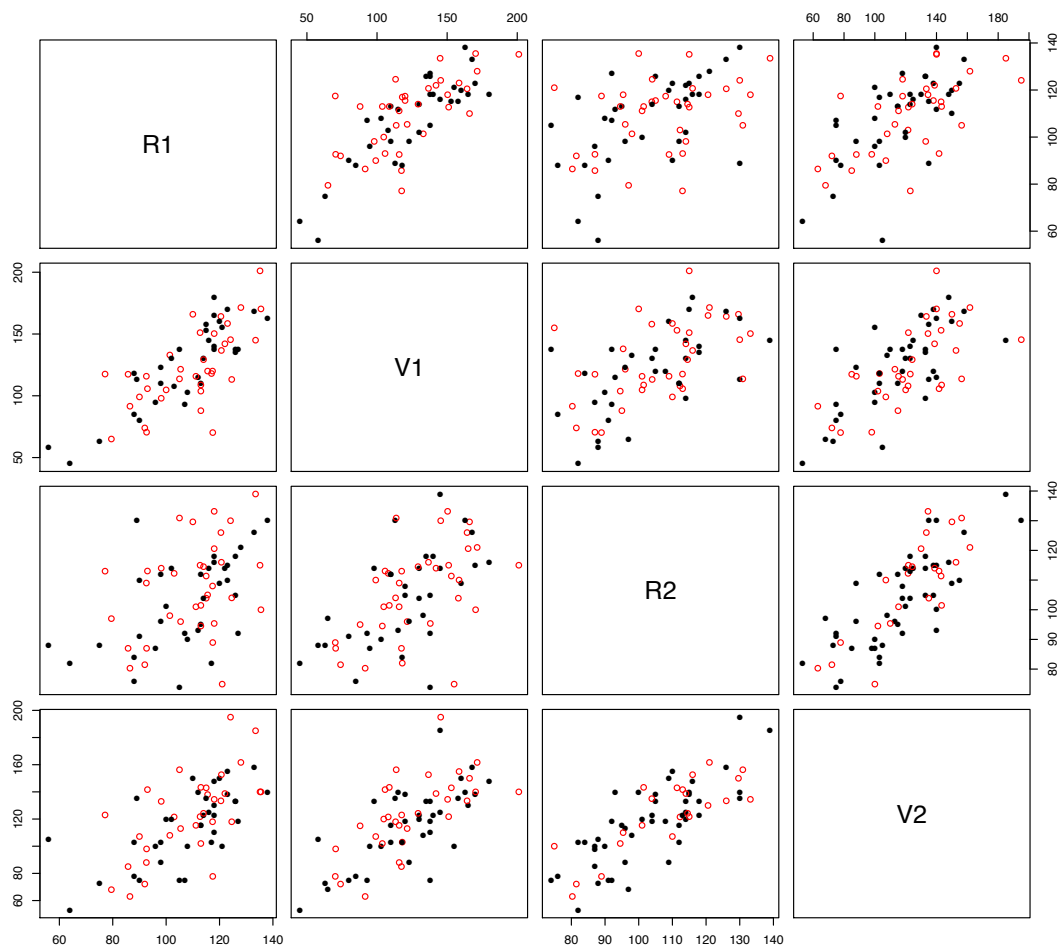|	
cdens	Component density (without mixing proportions).
dens	Mixture density.

Figure 32: Example of missing data imputations obtained via the `mix` package for the continuous variables in the `stlouis` dataset. The closed black circles correspond to nonmissing data, while the open red circles correspond to imputed missing values.

12.4   Model-based Clustering

mclustBIC	BIC computation; clusters and models through `summary`.
Mclust	Model-based clustering.
clustCombi	Combining Gaussian mixture components for clustering.

12.5   Model-based Density Estimation

densityMclust	Density estimation via `Mclust`.

12.6   Discriminant Analysis

Class Densities as Mixture Components:

cv1EMtrain	Training via leave-one-out crossvalidation.
bicEMtrain	Training via BIC.
estep	E-step of the EM algorithm.
mstep	M-step of the EM algorithm.
MclustDA	Model-based classification using `modelType = "EDDA"`.

Parameterized Gaussian Mixture for Class Densities (MclustDA):

MclustDA	Model-based classification using default argument `modelType = "MclustDA"`.
predict.MclustDA	MclustDA classification.
cv.MclustDA	k-fold cross-validation for MclustDA.

12.7   Bayesian Regularization

priorControl	specify a prior distribution.
defaultPrior	default prior distribution.

12.8   Support for Modeling and Classification

mclust.options	set or retrieve mclust options.
mclustModel	specify an mclust model.
mclustModelNames	description for an mclust model.
mclustVariance	template for (co)variance specification of an mclust model.
emControl	specify parameters affecting EM.
map	Convert conditional probabilities to a classification.
unmap	Convert a classification to indicator variables.
bic	BIC for parameterized Gaussian mixture models.
sim	Simulate data from a parameterized Gaussian mixture model.
mapClass	Mapping between two classifications.
classError	Classification error.
adjustedRandIndex	Adjusted Rand Index.
sigma2decomp	Convert mixture covariances to decomposition form.
decomp2sigma	Convert decomposition form to mixture covariances.
nVarParams	Number of variance parameters.
imputeData	Missing data imputation using the mix package.

12.9 Plotting Functions

Univariate Data

mclust1Dplot	Classification, uncertainty, density and/or classification errors.
plot.densityMclust	Plot method associated with the densityMclust function.

Bivariate Data

mclust2Dplot	Classification, uncertainty, and/or classification errors.
surfacePlot	Contour, image, or perspective plot of either density or uncertainty.

More than Two Dimensions

Classification, uncertainty, and/or classification errors:

coordProj	Coordinate projections.
randProj	Random projections.
MclustDR	Dimension reduction for model-based clustering and classification.

12.9.1 Other Plotting Functions

clPairs	Pairs plot showing classification.
imputePairs	Pairs plot to show missing data imputations.
uncerPlot	Relative uncertainty of misclassified observations.
plot.Mclust	Plots associated with Mclust results.
plot.mclustBIC	BIC plot associated with mclustBIC results.
plot.mclustDA	Plots associated with mclustDA results.
plot.mclustDR	Plots associated with mclustDR results.
plot.clustCombi	Plots associated with clustCombi results.

A Appendix

A.1 Clustering Models

mclust usually assumes a normal or Gaussian mixture model

$$\prod_{i=1}^{n}\sum_{k=1}^{G}\tau_k\ \phi_k(\mathbf{x}_i\mid\mu_k,\Sigma_k),$$

where where $\mathbf{x}$ represents the data, G is the number of components, τ_k is the probabilty that an observation belongs to the kth component $(\tau_k\geq 0;\quad\sum_{k=1}^{G}\tau_k=1)$, and

$$\phi_k(\mathbf{x}\mid\mu_k,\Sigma_k)=(2\pi)^{-\frac{p}{2}}|\Sigma_k|^{-\frac{1}{2}}\exp\left\{-\frac{1}{2}(\mathbf{x}_i-\mu_k)^T\Sigma_k^{-1}(\mathbf{x}_i-\mu_k)\right\}. \qquad (1)$$

The exception is for model-based hierarchical clustering, for which the model used is the classification likelihood with a parameterized normal distribution assumed for each class:

$$\prod_{i=1}^{n}\phi_{\ell_i}(\mathbf{x}_i\mid\mu_{\ell_i},\Sigma_{\ell_i}),$$

where the ℓ_i are labels indicating a unique classification of each observation: $\ell_i=k$ if $\mathbf{x}_i$ belongs to the kth component.

The components or clusters in both these models are ellipsoidal, centered at the means μ_k. The covariances Σ_k determine their other geometric features. Each covariance matrix is parameterized by eigenvalue decomposition in the form

$$\Sigma_k=\lambda_k D_k A_k D_k^T,$$

where D_k is the orthogonal matrix of eigenvectors, A_k is a diagonal matrix whose elements are proportional to the eigenvalues of Σ_k, and λ_k is a scalar (Banfield and Raftery 1993). The orientation of the principal components of Σ_k is determined by D_k, while A_k determines the shape of the density contours; λ_k specifies the volume of the corresponding ellipsoid, which is proportional to $\lambda_k^d|A_k|$, where d is the data dimension. Characteristics (orientation, volume and shape) of distributions are usually estimated from the data, and can be allowed to vary between clusters, or constrained to be the same for all clusters [27, 2, 8]. This parameterization includes but is not restricted to well-known variance models that are associated with various criterion for hierarchical clustering, such as equal-volume spherical variance $(\Sigma_k=\lambda I)$ for the sum of squares criterion [33], constant variance [20], and unconstrained variance [30].

In one dimension, there are just two models: E for equal variance and V for varying variance. In more than one dimension, the model identifiers code geometric characteristics of the model. For example, EVI denotes a model in which the volumes of all clusters are equal (E), the shapes of the clusters may vary (V), and the orientation is the identity (I). Clusters in this model have diagonal covariances with orientation parallel to the coordinate axes. Parameters associated with characteristics designated by E or V are determined from the data. Table 1 shows the various multivariate model options currently available in mclust for hierarchical clustering (denoted HC) and EM. These are a subset of the parameterizations discussed in [8], which gives details of the EM algorithm for maximum likelihood estimation for these models.

A.2 Modeling Noise and Outliers

`mclust` uses a mixture model which has a single term representing noise as a first order Poisson process to handle noisy data:

$$\prod_{i=1}^{n} \left[\frac{\tau_0}{V} + \sum_{k=1}^{G} \tau_k \phi_k(\mathbf{x}_i \mid \theta_k) \right], \tag{2}$$

in which V is the hypervolume of the data region, and $\tau_k \geq 0$; $\sum_{k=0}^{G} \tau_k = 1$. This model has been used successfully in a number of applications [2, 9, 6, 7].

The basic model-based clustering method needs to be modified when the data contains noise. First, a good initial noise estimate must be obtained. Some possible methods for denoising include a Voronoï method [1], a nearest-neighbor method [5], and robust covariance estimation [32]. The function `NNclean` in the contributed `R` package `prabclus` is an implementation of the nearest-neighbor method. The function `cov.nnve` in the contributed `R` package `covRobust` is an implementation of robust covariance estimation. Next, hierarchical clustering is applied to the denoised data. Finally, EM based on the Gaussian model with the added noise term (2) is applied to the entire data set, with the data removed in the denoising process as the initial noise estimate.

A.3 Model Selection via BIC

Several measures have been proposed for choosing the clustering model (parameterization and number of clusters); see, e.g., Chapter 6 of [26]. We use the Bayesian Information Criterion (BIC) approximation to the Bayes factor [29], which adds a penalty to the loglikelihood based on the number of parameters, and has performed well in a number of applications (e.g. [11, 13]). The BIC has the form

$$\text{BIC} \equiv 2 \, \text{loglik}_{\mathcal{M}}(\mathbf{x}, \theta_k^*) - (\# \text{ params})_{\mathcal{M}} \, \log(n), \tag{3}$$

where $\text{loglik}_{\mathcal{M}}(\mathbf{x}, \theta_k^*)$ is the maximized loglikelihood for the model and data, $(\# \text{ params})_{\mathcal{M}}$ is the number of independent parameters to be estimated in the model $\mathcal{M}$, and n is the number of observations in the data.

A.4 Adding a Prior to the Model

By default, `mclust` does not use a prior for modeling. However, users can optionally specify a conjugate prior of the type described in this section. For univariate data, we use a normal prior on the mean (conditional on the variance):

$$\begin{aligned} \mu \mid \sigma^2 \quad &\sim \quad \mathcal{N}(\mu_{\mathcal{P}}, \sigma^2/\kappa_{\mathcal{P}}) \\ &\propto \quad \left(\sigma^2\right)^{-\frac{1}{2}} \exp\left\{ -\frac{\kappa_{\mathcal{P}}}{2\sigma^2} \left(\mu - \mu_{\mathcal{P}}\right)^2 \right\} \end{aligned} \tag{4}$$

and an inverse gamma prior on the variance:

$$\begin{aligned} \sigma^2 \quad &\sim \quad \text{inverseGamma}(\nu_{\mathcal{P}}/2, \varsigma_{\mathcal{P}}^2/2) \\ &\propto \quad \left(\sigma^2\right)^{-\frac{\nu_{\mathcal{P}}+2}{2}} \exp\left\{ -\frac{\varsigma_{\mathcal{P}}^2}{2\sigma^2} \right\}. \end{aligned} \tag{5}$$

For multivariate data, we use a normal prior on the mean (conditional on the covariance matrix):

$$\begin{aligned} \mu \mid \Sigma \quad &\sim \quad \mathcal{N}(\mu_{\mathcal{P}}, \Sigma/\kappa_{\mathcal{P}}) \\ &\propto \quad |\Sigma|^{-\frac{1}{2}} \exp\left\{ -\frac{\kappa_{\mathcal{P}}}{2} \text{trace} \left[(\mu - \mu_{\mathcal{P}})^T \Sigma^{-1} (\mu - \mu_{\mathcal{P}}) \right] \right\}, \end{aligned} \tag{6}$$

and an inverse Wishart prior on the covariance matrix:

$$\Sigma \;\sim\; \text{inverseWishart}(\nu_{\mathcal{P}}, \Lambda_{\mathcal{P}})$$

$$\propto\;\; |\Sigma|^{-\frac{\nu_{\mathcal{P}}+d+1}{2}} \exp\left\{-\frac{1}{2}\text{trace}\left[\Sigma^{-1}\Lambda_{\mathcal{P}}^{-1}\right]\right\}. \tag{7}$$

The hyperparameters $\mu_{\mathcal{P}}$, $\kappa_{\mathcal{P}}$, and $\nu_{\mathcal{P}}$ are called the *mean*, *shrinkage* and *degrees of freedom*, respectively. Parameters $\varsigma_{\mathcal{P}}^2$ (a scalar) and $\Lambda_{\mathcal{P}}$ (a matrix) are the *scale* of the prior distribution in the univariate and multivariate cases, respectively. These priors are called *conjugate priors* for the normal distribution because the posterior can be expressed as the product of a normal distribution and an inverse gamma or Wishart distribution. With the prior, a modified version of the BIC, in which the MLE is replaced by the MAP, is used to choose the number of clusters. Details on model-based clustering with a prior can be found in [17].

Functions `priorControl` and `defaultPrior` are provided in `mclust` for specifying a prior. When called with defaults, the following choices are made for the prior hyperparameters:

$\mu_{\mathcal{P}} =$ the mean of the data.

$\kappa_{\mathcal{P}} =$ .01

The posterior mean $\dfrac{n_k \bar{y}_k + \kappa_{\mathcal{P}}\mu_{\mathcal{P}}}{\kappa_{\mathcal{P}} + n_k}$ can be viewed as adding $\kappa_{\mathcal{P}}$ observations with value $\mu_{\mathcal{P}}$ to each group in the data. The value we used was determined by experimentation; values close to and bigger than 1 caused large perturbations in the modeling in cases where there were no missing BIC values without the prior. The value .01 resulted in BIC curves that appeared to be smooth extensions of their counterparts without the prior.

$\nu_{\mathcal{P}} = d + 2$

Analogously to the univariate case, the marginal prior distribution of μ is a t distribution centered at $\mu_{\mathcal{P}}$ with $\nu_{\mathcal{P}} - d + 1$ degrees of freedom. The mean of this distribution is $\mu_{\mathcal{P}}$ provided that $\nu_{\mathcal{P}} > d$, and it has a finite covariance matrix provided $\nu_{\mathcal{P}} > d + 1$ (see, e. g. Schafer 1997). We chose the smallest integer value for the degrees of freedom that gives a finite covariance matrix.

$\varsigma_{\mathcal{P}}^2 = \dfrac{\text{sum}(\text{diag}(\text{var}(\text{data})))/d}{G^{2/d}}$ (for univariate models, and multivariate spherical or diagonal models)

The average of the diagonal elements of the empirical covariance matrix of the data divided by the square of the number of components to the $1/d$ power. This is roughly equivalent to partitioning the range of the data into G intervals of fairly equal size.

$\Lambda_{\mathcal{P}} = \dfrac{\text{var}(\text{data})}{G^{2/d}}$ (for multivariate ellipsoidal models)

The empirical covariance matrix of the data divided by the square of the number of components to the $1/d$ power.

References

[1] D. Allard and C. Fraley. Nonparametric maximum likelihood estimation of features in spatial point processes using Voronoï tessellation. *Journal of the American Statistical Association*, 92:1485–1493, 1997.

[2] J. D. Banfield and A. E. Raftery. Model-based Gaussian and non-Gaussian clustering. *Biometrics*, 49:803–821, 1993.

[3] J.-P. Baudry, A. E. Raftery, G. Celeux, K. Lo and R. Gottardo. Combining mixture components for clustering. *Journal of Computational and Graphical Statistics*, 19(2):332-353, 2010.

[4] H. Bensmail, and G. Celeux. Regularized Gaussian Discriminant Analysis Through Eigenvalue Decomposition. *Journal of the American Statistical Association*, 91:1743–1748, 1996.

[5] S. D. Byers and A. E. Raftery. Nearest neighbor clutter removal for estimating features in spatial point processes. *Journal of the American Statistical Association*, 93:577–584, 1998.

[6] J. G. Campbell, C. Fraley, F. Murtagh, and A. E. Raftery. Linear flaw detection in woven textiles using model-based clustering. *Pattern Recognition Letters*, 18:1539–1548, 1997.

[7] J. G. Campbell, C. Fraley, D. Stanford, F. Murtagh, and A. E. Raftery. Model-based methods for real-time textile fault detection. *International Journal of Imaging Systems and Technology*, 10:339–346, 1999.

[8] G. Celeux and G. Govaert. Gaussian parsimonious clustering models. *Pattern Recognition*, 28:781–793, 1995.

[9] A. Dasgupta and A. E. Raftery. Detecting features in spatial point processes with clutter via model-based clustering. *Journal of the American Statistical Association*, 93:294–302, 1998.

[10] C. Fraley. Algorithms for model-based Gaussian hierarchical clustering. *SIAM Journal on Scientific Computing*, 20:270–281, 1998.

[11] C. Fraley and A. E. Raftery. How many clusters? Which clustering method? - Answers via model-based cluster analysis. *The Computer Journal*, 41:578–588, 1998.

[12] C. Fraley and A. E. Raftery. MCLUST: Software for model-based cluster analysis. *Journal of Classification*, 16:297–306, 1999.

[13] C. Fraley and A. E. Raftery. Model-based clustering, discriminant analysis and density estimation. *Journal of the American Statistical Association*, 97:611–631, 2002.

[14] C. Fraley and A. E. Raftery. Enhanced software for model-based clustering, density estimation, and discriminant analysis: MCLUST. *Journal of Classification*, 20:263–286, 2003.

[15] C. Fraley and A. E. Raftery. Model-based microarray image analysis. *R News*, 6:60–63, 2006.

[16] C. Fraley and A. E. Raftery. Some applications of model-based clustering in chemistry. *R News*, 6:17–23, 2006.

[17] C. Fraley and A. E. Raftery. Bayesian regularization for normal mixture estimation and model-based clustering. *Journal of Classification*, 24:155–181, 2007.

[18] C. Fraley and A. E. Raftery. Model-based methods of classification: using the `mclust` software in chemometrics. *Journal of Statistical Software*, 18(6), January 2007.

[19] C. Fraley, A. E. Raftery, and R. Wehrens. Incremental model-based clustering for large datasets with small clusters. *Journal of Computational and Graphical Statistics*, 14:1–18, 2005.

[20] H. P. Friedman and J. Rubin. On some invariant criteria for grouping data. *Journal of the American Statistical Association*, 62:1159–1178, 1967.

[21] G. H. Golub and C. F. Van Loan. *Matrix Computations*. Johns Hopkins, 3rd edition, 1996.

[22] L. Hubert and P. Arabie. Comparing partitions. *Journal of Classification*, 2:193–218, 1985.

[23] R. J. A. Little and D. B. Rubin. *Statistical Analysis with Missing Data*. Wiley, 2nd edition, 2002.

[24] C. Loader. *Local Regression and Likelihood*. Springer, New York, 1999.

[25] K. V. Mardia, J. T. Kent, and J. M. Bibby. *Multivariate Analysis*. Academic Press, 1979.

[26] G. J. McLachlan and D. Peel. *Finite Mixture Models*. Wiley, 2000.

[27] F. Murtagh and A. E. Raftery. Fitting straight lines to point patterns. *Pattern Recognition*, 17:479–483, 1984.

[28] J. L. Schafer. *Analysis of Incomplete Multivariate Data by Simulation*. Chapman and Hall, 1997.

[29] G. Schwarz. Estimating the dimension of a model. *The Annals of Statistics*, 6:461–464, 1978.

[30] A. J. Scott and M. J. Symons. Clustering methods based on likelihood ratio criteria. *Biometrics*, 27:387–397, 1971.

[31] L. Scrucca. Dimension reduction for model-based clustering. *Statistics and Computing*, 20(4):471–484, 2010.

[32] N. Wang and A. E. Raftery. Nearest neighbor variance estimation (NNVE): Robust covariance estimation via nearest neighbor cleaning (with discussion). *Journal of the American Statistical Association*, 97:994–1019, 2002.

[33] J. H. Ward. Hierarchical groupings to optimize an objective function. *Journal of the American Statistical Association*, 58:234–244, 1963.

[34] R. Wehrens, L. Buydens, C. Fraley, and A. Raftery. Model-based clustering for image segmentation and large datasets via sampling. *Journal of Classification*, 21:231–253, 2004.